

H12

# DBMS Architectures

Fundamentals of Databases

Alvaro A A Fernandes, SCS, UoM

[COMP23111 2014-2015 Lecture 12 of 12]

# Acknowledgements

- These slides are adaptations (mostly minor, but some major) of material authored and made available to instructors by **A. Silberschatz, H. F. Korth, and S. Sudarshan** to accompany their textbook **Database System Concepts, 6th Edition, McGraw-Hill, 2006, 978-0-07-352332-3**.
- Copyright remains with them, whom I thank.
- Any errors are my responsibility.

# In Previous Lectures

- We learned that data is an enterprise asset of such importance that specialized software systems, called DBMSs, are crucial to manage it well and that using files alone won't do.
- We learned the importance of separating concerns and of adopting different levels of abstraction in designing and implementing databases.
- We learned how data models lead to schemas and instances that enable a logical view of the data.
- We learned how to capture data requirements in the form of an (E)ER conceptual model.
- We learned how to derived from this conceptual model a logical one in the relational case, which is directly implementable in a DBMS.
- We learned how SQL has been extended to Turing-completeness and how views and trigger enrich our capability to capture and enforce domain semantics.
- We learned how DBMSs use the notion of transactions to manage concurrency and recovery.
- We learned how file organizations and indexed enable a logical-to-physical abstraction level at the storage layer too.


# In This Lecture

- What is the internal architecture of a classical DBMS like?
  - ▶ What are the main subsystems?
  - ▶ What are the main classes of users of DBMSs?
- What are the main kinds of database architecture in use today?
  - ▶ How are client-server designs realized in DBMS architecture?
  - ▶ What kinds of server are there?
  - ▶ What parallelization and distribution approaches are there for DBMSs?
  - ▶ How do DBMSs scale?

What is the internal architecture of a classical DBMS like? What are the main subsystems? What are the main classes of users?

# DBMS Architecture: Components/Interactions


7



# Main DBMS Subsystems :: Storage Management

8

- Storage manager is a program module that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system.
- The storage manager is responsible to the following tasks:
  - ▶ Interaction with the file manager
  - ▶ Efficient storing, retrieving and updating of data
- Design/implementation issues:
  - ▶ Storage access
  - ▶ File organization
  - ▶ Indexing and hashing

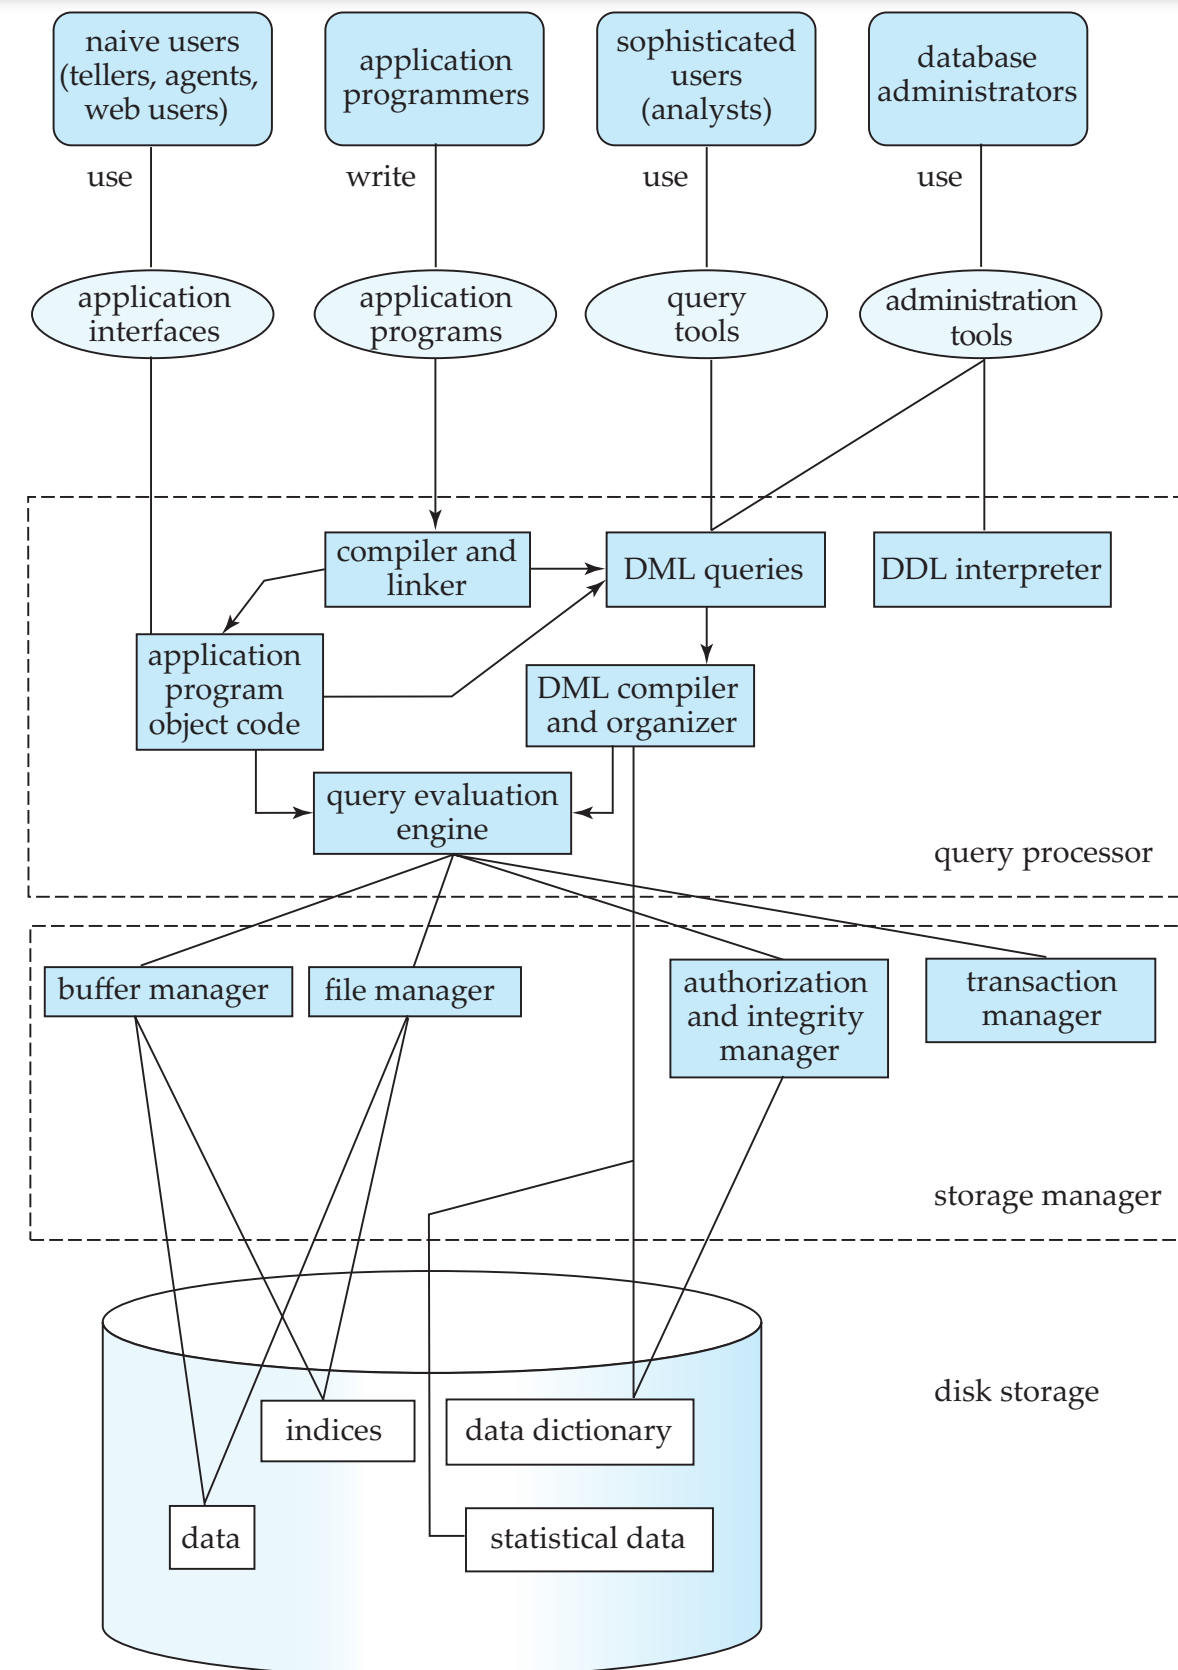




# Main DBMS Subsystems :: Transaction Management

9


- What if the system fails?
- A transaction is a collection of operations that performs a single logical function in a database application
- Transaction-management component ensures that the database remains in a consistent (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures.



# Main DBMS Subsystems :: Concurrency Controller

10

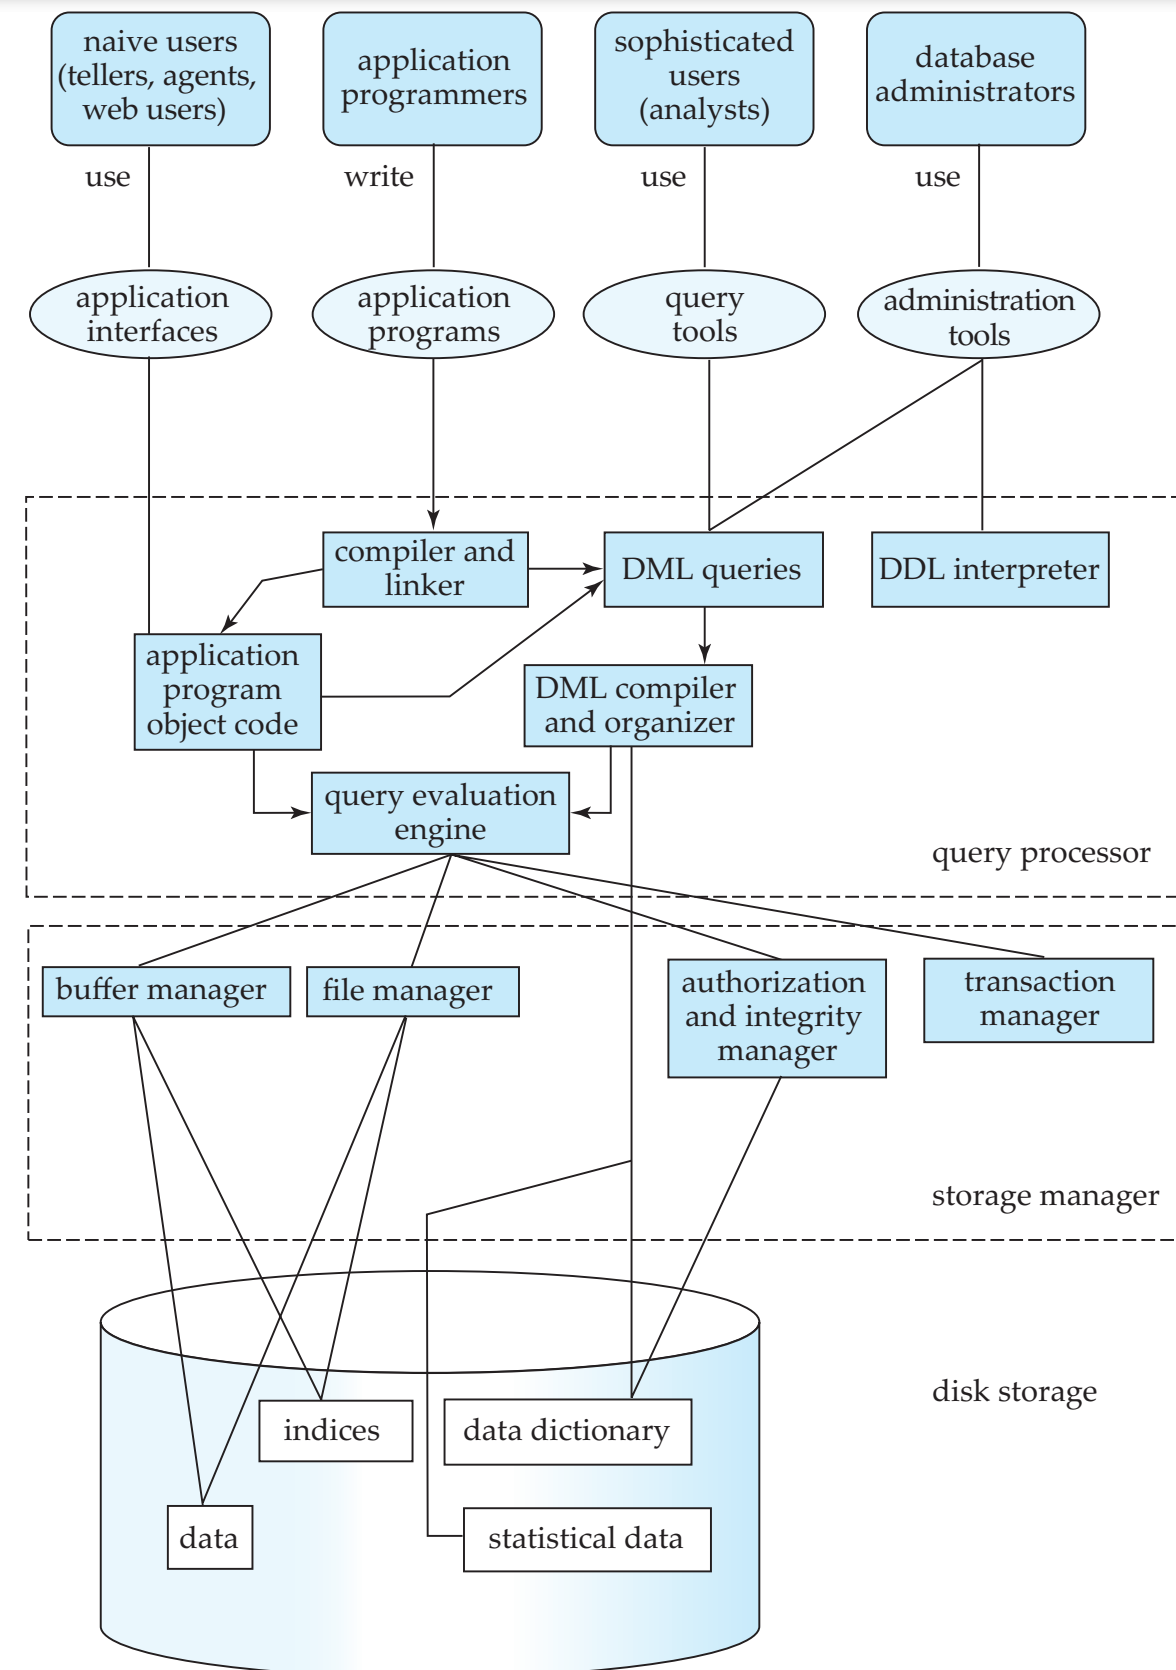
- What if more than one user is concurrently updating the same data?
- Concurrency controller ensures that the interaction among the concurrent transactions do not compromise the consistency of the database.



# Main DBMS Subsystems: Query Processor

11

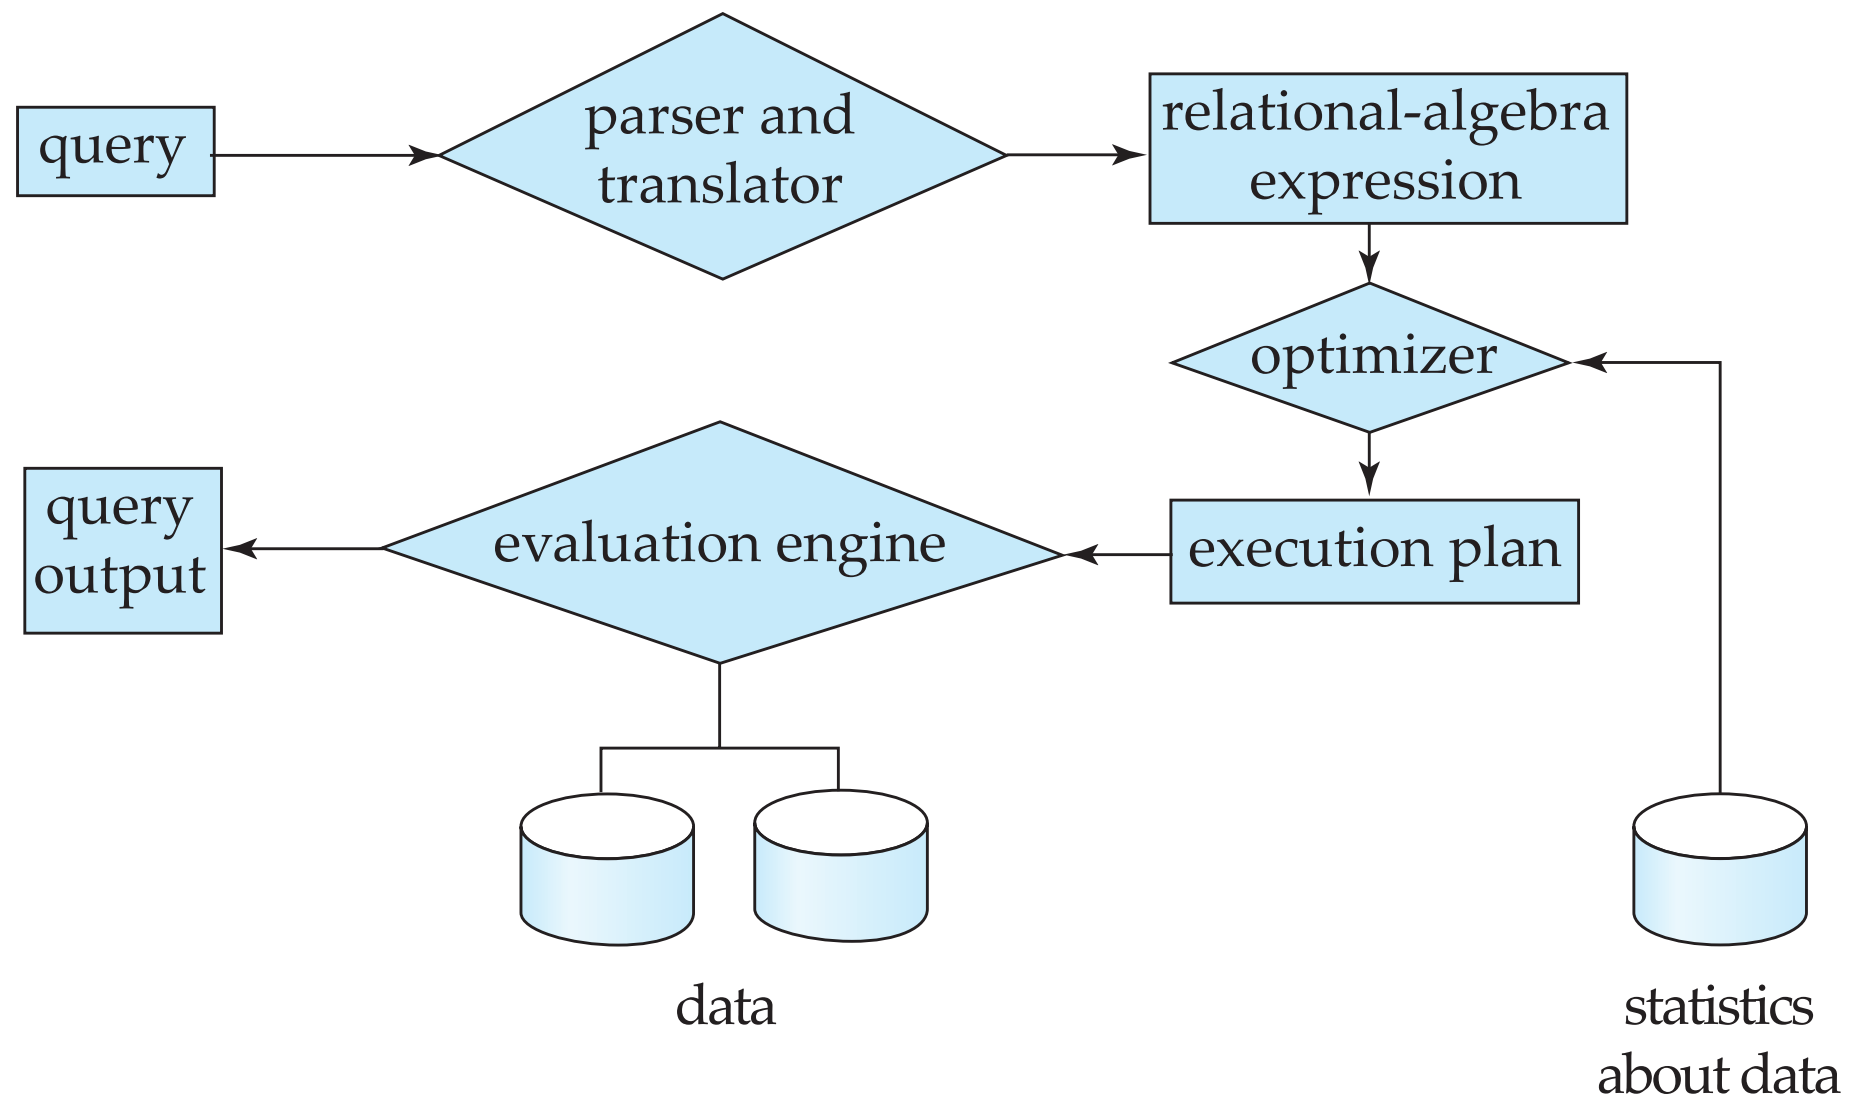
- Here, by query is meant DDL, QL and DML expressions
- Code in the database languages is compiled and installed in the DBMS.
- Application programs in general-purpose languages can invoke compiled queries or pass queries for evaluation using the libraries and APIs exposed by the DBMS
- The HCI interfaces invoke bits of application that interact with the DBMS.



# Main Phases of Query Processing

12


- Parsing and translation
- Optimization
- Evaluation



# Inside the Query Processor: Query Optimizer

13

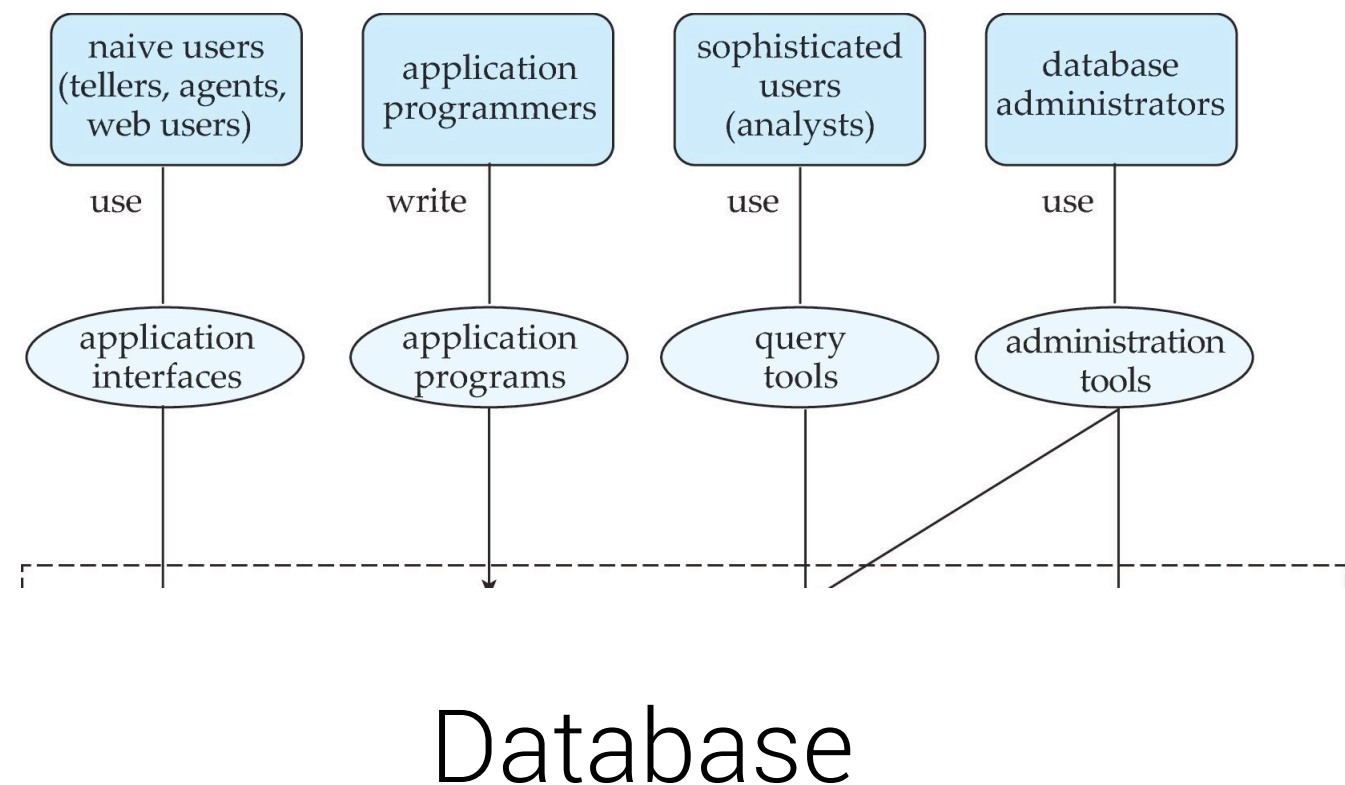
- There are alternative ways of evaluating a given query
  - ▶ Many equivalent logical expressions
  - ▶ Many algorithms for each logical operation with different associated costs
- Cost difference between a good and a bad way of evaluating a query can be enormous
- Need to estimate the cost of operations
  - ▶ Depends critically on statistical information about relations which the database must maintain
  - ▶ Need to estimate statistics for intermediate results to compute cost of complex expressions



# DBMS Users: Roles

14

- Database administrators assign privileges to users and monitor performance.
- They continually study the statistics and revise configuration parameters to fine-tune performance.
- Sophisticated users (e.g., database designers) use CLI to implement and maintain logical models.
- Application programmers write software that uses APIs over the database, including HCI interfaces for naïve users.
- Naïve users rely on HCI interfaces for both back-room and customer-facing tasks.



What are the main kinds of database architecture in use today?

# Database Architecture: Centralized to Distributed

16

- The architecture of a database systems is greatly influenced by the underlying computer system on which the database is running:
  - ▶ Centralized
  - ▶ Client-server
  - ▶ Parallel (multi-processor)
  - ▶ Distributed



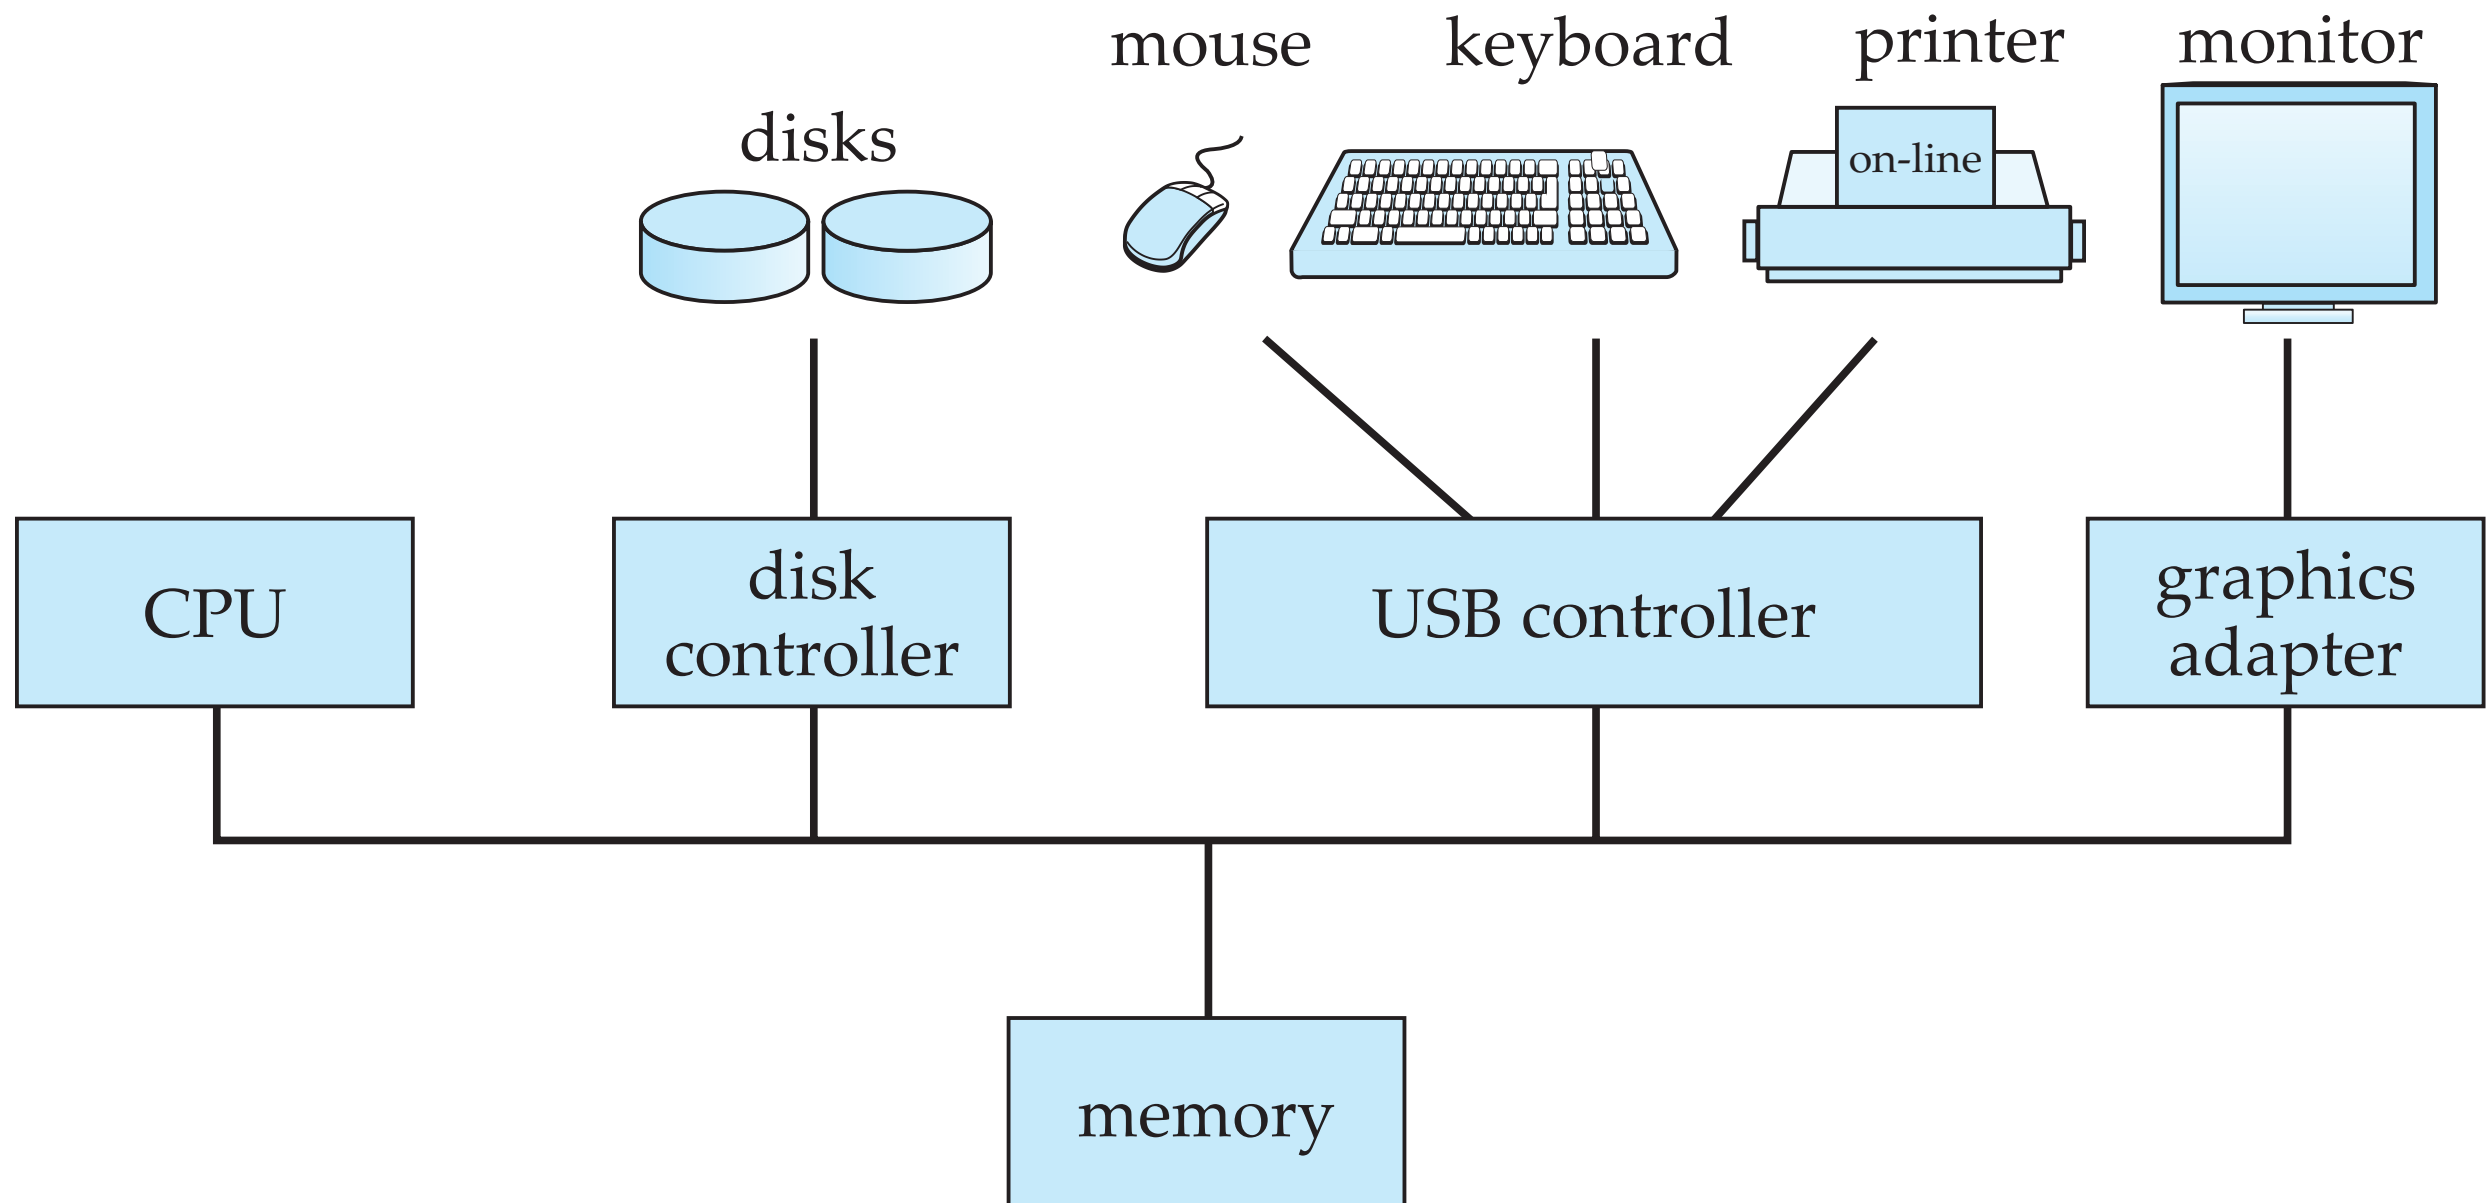
# Centralized Systems

17

- Run on a single computer system and do not need to interact with other computer systems.
- General-purpose computer systems: one to a few CPUs and a number of device controllers that are connected through a common bus that provides access to shared memory
- Single-user systems (e.g., personal computer or workstation): desk-top unit, logically for single user, logically one CPU and one secondary mass storage unit; the OS need not usually worry about more than one user at a time
- Multi-user systems: logically and physically many disks, much more memory, multiple CPUs, and a multi-user OS to serve a large number of users who are connected via terminals, workstations, PCs

# Centralized Systems


18



# Client-Server Systems

# Client-Server Systems

20

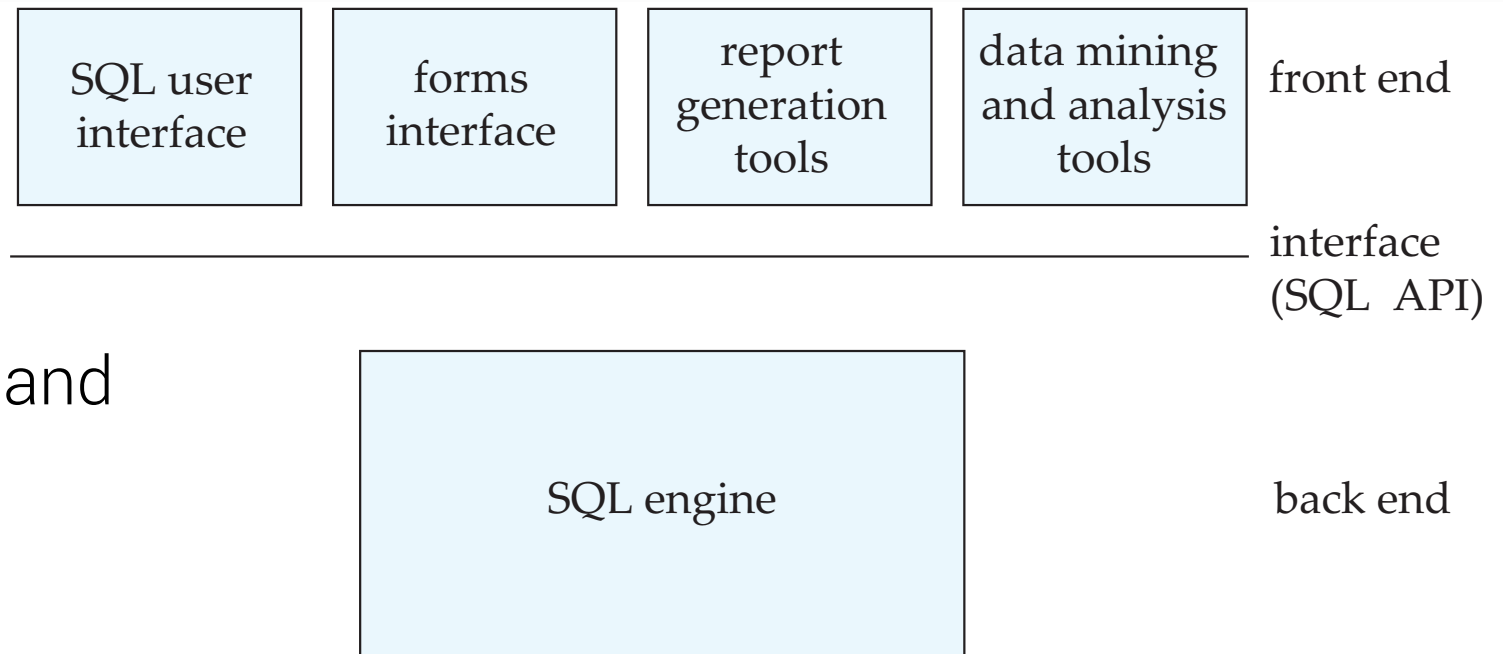


- Server systems are large multi-user systems that aim to satisfy requests generated at many client systems, each (logically) a single-user system.
- There is an asymmetry assumption that the server is many times better resourced than any client.

# Client-Server Systems

21

- Database functionality can be divided into:



- ▶ Back-end: manages access structures, query evaluation and optimization, concurrency control and recovery;
  - ▶ Front-end: consists of tools such as forms, report-writers, and graphical user interface facilities.
- The interface between the front-end and the back-end is
    - ▶ directly through a SQL engine (i.e., a query processor)
    - ▶ indirectly through an API.

# Client-Server Systems

- Some advantages of replacing mainframes (i.e., very large centralized computer systems) with networks of workstations or personal computers (i.e., single-user systems) acting as clients by way of a connection to back-end server machines (i.e., multi-user systems) are:
  - ▶ more functionality for the same cost,
  - ▶ flexibility in locating resources and expanding facilities,
  - ▶ better user interfaces,
  - ▶ easier maintenance.

# Types of Server


- Server systems can be broadly categorized into two kinds:
  - ▶ transaction servers, which are widely used in relational database systems
  - ▶ data servers, which are often used in object-based stores

# Transaction Servers



# Transaction Servers :: Communication


25



- Also called query server or SQL server systems
- ▶ Clients send requests to the server
- ▶ Transactions are executed at the server
- ▶ Results are shipped back to the client.

# Transaction Servers :: Mechanisms

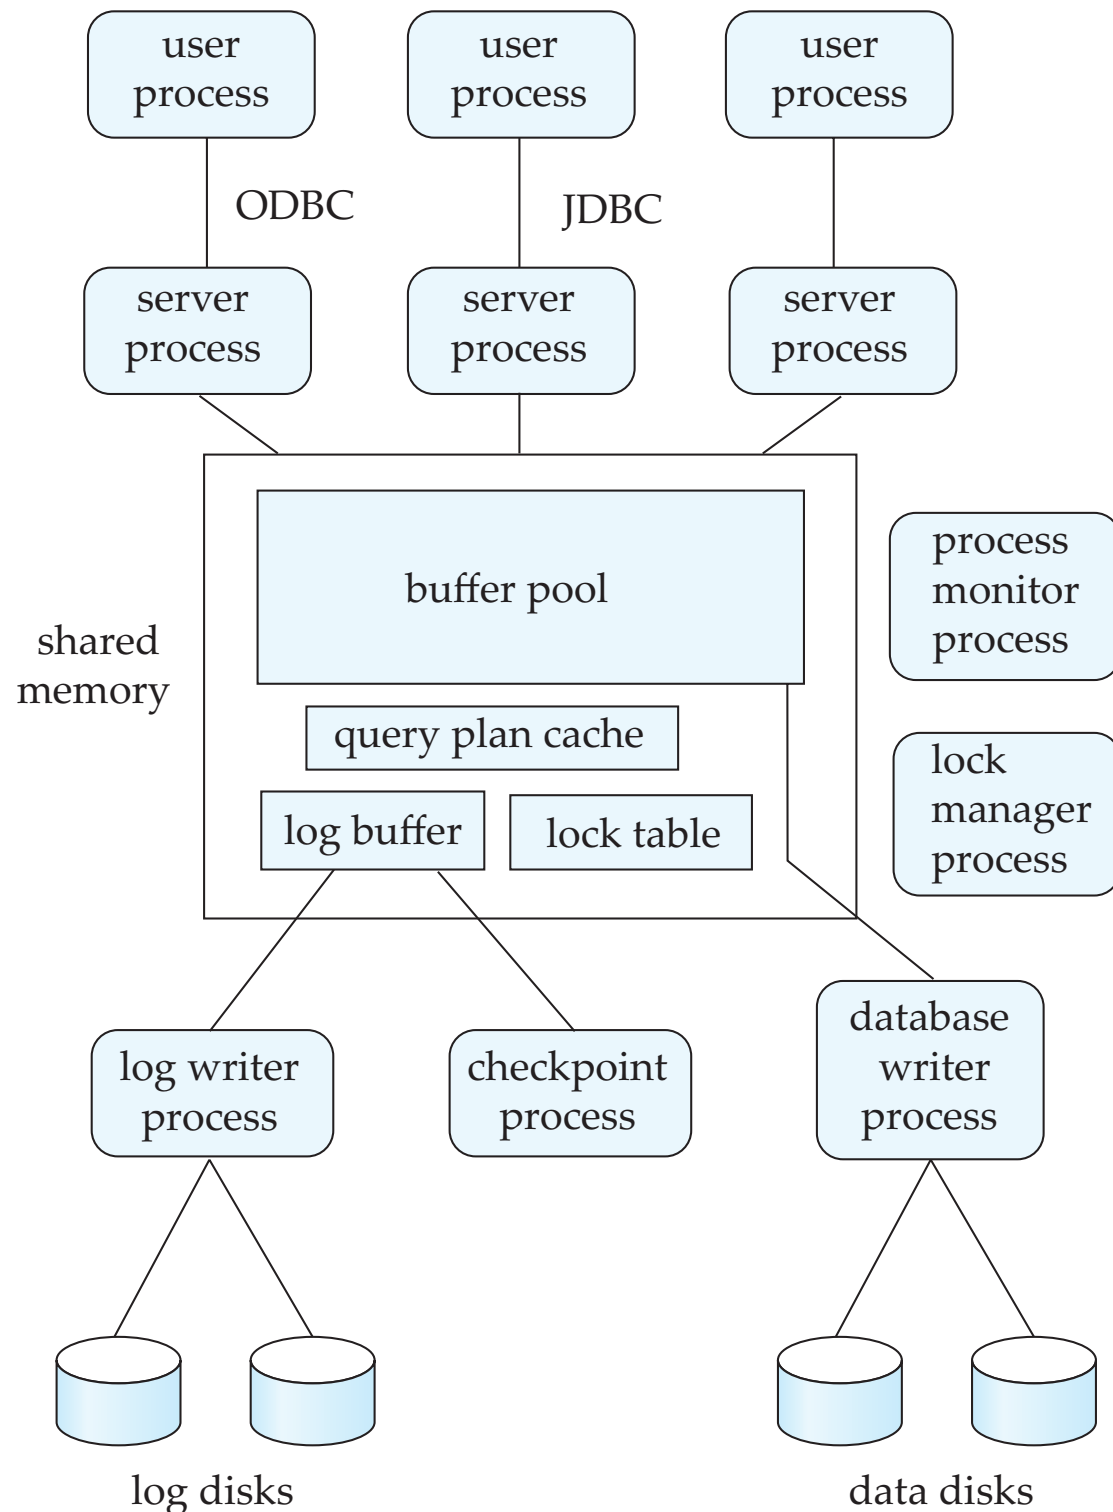
26



- Requests are specified in SQL, and communicated to the server through a remote procedure call (RPC) mechanism.
- Transactional RPC allows many RPC calls to form a transaction.
- Open Database Connectivity (ODBC) is a C-language API for connecting to a server, sending SQL requests, and receiving results.
- JDBC is Java-based ODBC.

# Transaction Servers :: Querying Processes


27



- In a transaction server, many processes cooperate to implement the database architecture we saw earlier.
- Queries are handled by server processes (i.e., these instantiate query processing functionality)
- These are typically multithreaded, typically massively so

# Transaction Servers :: Management Processes

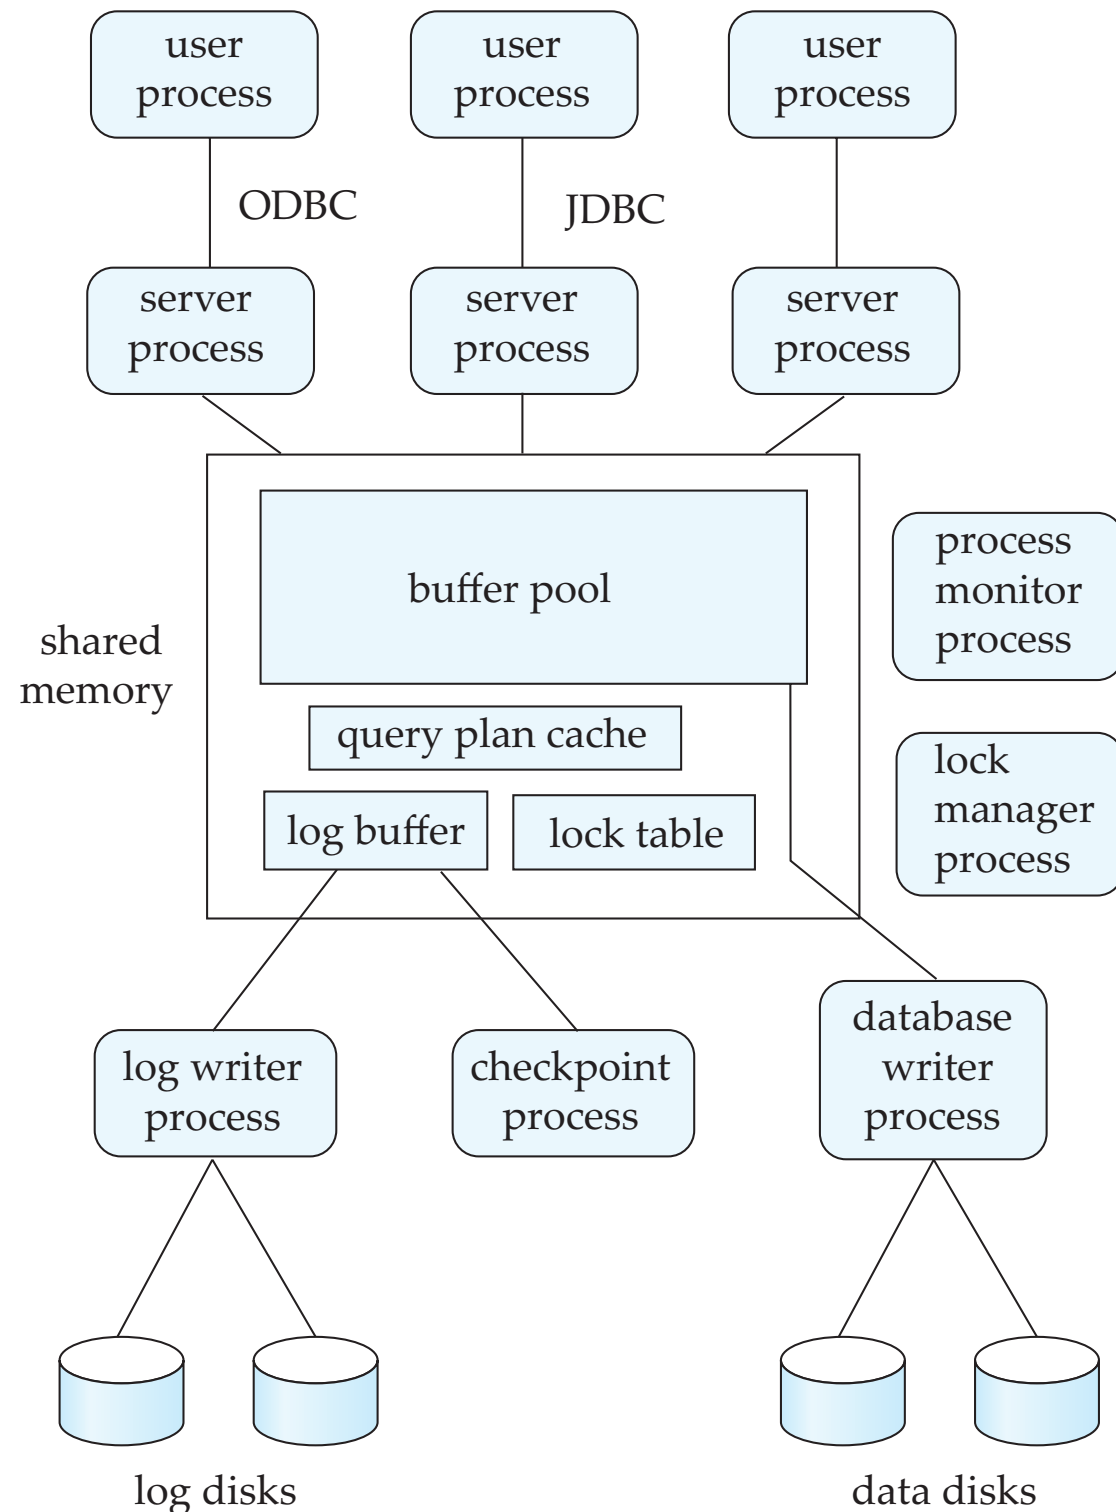
28



- Other processes control access to shared memory:
  - ▶ lock manager for concurrent access
  - ▶ database writer for traffic between volatile and persistent storage
  - ▶ process monitor, checkpoint and log-writer for transactional semantics (e.g., commits and aborts/roll-backs, recovery, etc.)
- There are many trade-offs at play here: all the helper processes (i.e., lock manager, buffer writers/off-loaders, etc.) ensure semantic correctness but risk being performance drains

# Transaction Servers :: Data Structure Management

29



- Shared memory contains shared data, i.e.:
  - ▶ Buffer pool
  - ▶ Lock table
  - ▶ Log buffer
  - ▶ Cached query plans (reused if same query submitted again)
- All database processes can access shared memory
- To ensure that no two processes are accessing the same data structure at the same time, databases systems implement some form of mutual exclusion

# Data Servers

# Data Servers, Classically

31

- Used in high-speed LANs, in cases where
  - ▶ The clients are comparable in processing power to the server
  - ▶ The tasks to be executed are compute-intensive.
- Data is shipped to clients for processing, and then shipped results back to the server.
- This architecture requires full back-end functionality at the clients.
- Often used in object-based systems

# Data Servers, Recently

32

- More recently, so-called NoSQL systems focus on massive replication of essentially equipotent systems
- Specialization leads to data servers for, say, images or sound, to be separate from classical retrieval
- Note that in the NoSQL world, there are no queries in the sense we understand them here.



# Parallel Systems

- Parallel database systems consist of multiple processors and multiple disks connected by a fast interconnection network.
- A coarse-grain-parallel machine consists of a small number of powerful processors
- A massively-parallel or fine-grain-parallel machine utilizes thousands of smaller processors.

- Two main performance measures:
  - ▶ Throughput: the number of tasks that can be completed in a given time interval
  - ▶ Response time: the amount of time it takes to complete a single task from the time it is submitted
- The goal is to improve performance as we increase the degree of parallelism (i.e., the number of resources deployed in parallel)


$$\begin{aligned} |P| &= p \\ |S| &= 3 \end{aligned}$$

$$\begin{aligned} |P| &= p \\ |L| &= 4 \cdot |S| \end{aligned}$$

- Assume a problem of a given size  $|P| = p$  that is executing on a small system of size  $|S|$  is given to a larger system  $|L| = n \cdot |S|$
- The speed-up goal is to then solve  $P$  in time inversely proportional to  $n$ , i.e., the growth in resources
- The problem size remains constant

# Speed-Up

37



$ P  = p$	$ P  = p$
$ S  = 3$	$ L  = 4 \cdot  S $
$T_S = 1$	$T_L = 0.25$
$T_S/T_L = 4 \rightarrow \text{linear}$	

- Speed-up is defined as

$$\text{SpUP} = T_S/T_L$$

i.e., the elapsed-time  $T_S$  to solve  $P$  in  $S$  divided by the elapsed-time  $T_L$  to solve  $P$  in  $L$  where  $|L| = n \cdot |S|$ ,  $n > 1$

- Speed-up is linear if  $\text{SpUP} = n$
- Sub-linear speed-up means that the full investment on more resources is not fully recouped in elapsed-time reduction

# Scale-Up


38

$$\begin{array}{ll} |P| = p & |P'| = 4.p \\ |S| = 3 & |L| = 4.|S| \end{array}$$

- Assume a problem of a given size  $|P|$  that is executing on a small system of size  $|S|$ .
- Now, assume that faced with an increase of  $k$ ,  $k > 1$  in the size of the problem, we increase the system  $k$ -times
- The goal here is to solve the larger problem in the same time as before.
- Both the problem size and the system size grow.

# Scale-Up

39



- Scale-up is defined as

$$\text{ScUP} = T_S / T_L$$

i. e., the elapsed-time  $T_S$  to solve  $P$  in  $S$  divided by the elapsed-time  $T_L$  to solve  $P'$  in  $L$  where  $|P'| = k \cdot |P|$  and  $|L| = k \cdot |S|$ ,  $k > 1$

- Scale-up is linear if  $\text{ScUP} = 1$
- Sub-linear scale-up means that the full investment on more resources is not fully recouped in elapsed-time stability

$$\begin{array}{ll} |P| = p & |P'| = 4 \cdot p \\ |S| = 3 & |L| = 4 \cdot |S| \\ T_S = 2 & T_L = 2 \\ T_S / T_L = 1 & \rightarrow \text{linear} \end{array}$$

- A single large job; typical of most decision support queries and scientific simulation.
- Use an N-times larger computer on N-times larger problem.



# Transaction Scale-Up

41

- Numerous small queries submitted by independent users to a shared database; typical transaction processing and timesharing systems.
- N-times as many users submitting requests (hence, N-times as many requests) to an N-times larger database, on an N-times larger computer.
- Well-suited to parallel execution.

# Speed-Up and Scale-Up Limiting Factors

42

- Speed-up and scale-up are often sublinear due to:
  - ▶ Start-up costs: Cost of starting up multiple processes may dominate computation time, if the degree of parallelism is high.
  - ▶ Interference: Processes accessing shared resources (e.g., system bus, disks, or locks) compete with each other, thus spending time waiting on other processes, rather than performing useful work.
  - ▶ Skew: Increasing the degree of parallelism increases the variance in service times of parallel-executing tasks. Overall execution time determined by slowest of parallel-executing tasks.


# Parallel Database Architectures

43


- Shared memory -- processors share a common memory
- Shared disk -- processors share a common disk (array)
- Shared nothing -- processors share neither a common memory nor common disk (array)
- Hierarchical -- hybrid of the above architectures

# Parallel Database Architectures


44




(a) shared memory



(b) shared disk



(c) shared nothing



(d) hierarchical

# Shared-Memory Architectures

45

- Processors and disks have access to a common memory, typically via a bus or through an interconnection network.
- Extremely efficient communication between processors — data in shared memory can be accessed by any processor without having to move it using software.

# Shared-Memory Architectures

- Downside – architecture is not scalable beyond 32 or 64 processors since the bus or the interconnection network becomes a bottleneck
- Widely used for lower degrees of parallelism (4 to 8).

# Shared-Disk Architectures

47

- All processors can directly access all disks via an interconnection network, but the processors have private memories.
- The memory bus is not a bottleneck
- Architecture provides a degree of fault-tolerance — if a processor fails, the other processors can take over its tasks since the database is resident on disks that are accessible from all processors.

# Shared-Disk Architectures

- Downside: bottleneck now occurs at interconnection to the disk subsystem.
- Shared-disk systems can scale to a somewhat larger number of processors, but communication between processors is slower.
- Examples: IBM Sysplex and DEC clusters (now part of Compaq) running Rdb (now Oracle Rdb) were early commercial users



# Shared-Nothing Architectures

49

- Node consists of a processor, memory, and one or more disks. Processors at one node communicate with another processor at another node using an interconnection network. A node functions as the server for the data on the disk or disks the node owns.
- Data accessed from local disks (and local memory accesses) do not pass through interconnection network, thereby minimizing the interference of resource sharing.

# Shared-Nothing Architectures

50

- Shared-nothing multiprocessors can be scaled up to thousands of processors without interference.
- Main drawback: cost of communication and non-local disk access; sending data involves software interaction at both ends.
- Examples: Teradata, Tandem, Oracle-n CUBE

# Hierarchical Architectures

51

- Combines characteristics of shared-memory, shared-disk, and shared-nothing architectures.
- Top level is a shared-nothing architecture – nodes connected by an interconnection network, and do not share disks or memory with each other.
- Each node of the system could be a shared-memory system with a few processors.

# Hierarchical Architectures

52


- Alternatively, each node could be a shared-disk system, and each of the systems sharing a set of disks could be a shared-memory system.
- Reduce the complexity of programming such systems by distributed virtual-memory architectures
- Also called non-uniform memory architecture (NUMA)

# Distributed Systems

# Distributed Systems

54

- Data spread over multiple machines (also referred to as sites or nodes).
- Network interconnects the machines
- Data shared by users on multiple machines



# Distributed Databases :: Homogeneous

55

- Same software/schema on all sites, data may be partitioned among sites
- Goal: provide a view of a single database, hiding details of distribution

# Distributed Databases :: Heterogeneous

56

- Different software/schema on different sites
- Goal: integrate existing databases to provide useful functionality



# Distributed Databases :: Local v. Global

57

- A local transaction accesses data in the single site at which the transaction was initiated.
- A global transaction either accesses data in a site different from the one at which the transaction was initiated or accesses data in several different sites.

# Trade-Offs in Distributed Systems

58

- Sharing data – users at one site able to access the data residing at some other sites.
- Autonomy – each site is able to retain a degree of control over data stored locally.
- Higher system availability through redundancy – data can be replicated at remote sites, and system can function even if a site fails.

# Trade-Offs in Distributed Systems

59

- Disadvantage: added complexity required to ensure proper coordination among sites.
  - ▶ Software development cost.
  - ▶ Greater potential for bugs.
  - ▶ Increased processing overhead.

# Implementation Issues in Distributed Databases

60

- Atomicity needed even for transactions that update data at multiple sites
- The two-phase commit protocol (2PC) is used to ensure atomicity
  - ▶ Basic idea: each site executes transaction until just before commit, and then leaves final decision to a coordinator
  - ▶ Each site must follow decision of coordinator, even if there is a failure while waiting for coordinator's decision

# Implementation Issues in Distributed Databases


61

- 2PC is not always appropriate: other transaction models based on persistent messaging, and workflows, are also used
- Distributed concurrency control (and deadlock detection) required
- Data items may be replicated to improve data availability
- (More details later in this course unit)

# Types of Network

62


- Local-area networks (LANs) – composed of processors that are distributed over small geographical areas, such as a single building or a few adjacent buildings.
- Wide-area networks (WANs) – composed of processors distributed over a large geographical area.



# Types of Network

63

- WANs with continuous connection (e.g., the Internet) are needed for implementing distributed database systems
- Groupware applications such as Lotus notes can work on WANs with discontinuous connection:
  - ▶ Data is replicated.
  - ▶ Updates are propagated to replicas periodically.
  - ▶ Copies of data may be updated independently.
  - ▶ Non-serializable executions can thus result. Resolution is application dependent.



# Summary



# DBMS Components, Subsystems, Users, and Large-Scale Deployment Architectures

65

- The main components of DBMSs are the query processor and the storage manager.
- The query processor (QP) parses, translates and optimizes a query to generate an evaluation plan, which when executed by the query evaluation engine, produces the query results.
- In so doing the QP relies on the storage manager to regulate access to buffers in primary memory and to manage its movement from and to secondary memory.
- Important subsystems also include the concurrency and the transaction managers.
- The main classes of users of DBMSs are causal/occasional users, sophisticated users, application programmers and administrators.
- They rely on tools and APIs provided/exposed by the DBMS.
- The main kinds of database architectures used today are client-server, parallel and distributed.
- Each of these have variants (e.g., finely- or coarsely-grained parallel, WAN- or LAN-distributed).

# The End

- Thanks for attending the lectures.
- I hope you have enjoyed the course.