



# COMP25111

Operating Systems

55, 47, 41 mins

Lectures 11  
Virtual Memory (1)


<http://www.cs.man.ac.uk/~neville/COMP25111/Lecture11full/Lecture11full.html>

# COMP25111

## Operating Systems

### Lectures 11

### Virtual Memory (1)



© Copyright Richard Neville 2007

Dr Richard Neville

[r.neville@manchester.ac.uk](mailto:r.neville@manchester.ac.uk)



Room: G12 Kilburn Building, Bottom floor

Week

**7**

NOTE: The up-to-date version of this lecture is kept on the associated web site – available [on-line] @ Blackboard select: COMP25111 Introduction to Computer Systems [www.manchester.ac.uk/portal](http://www.manchester.ac.uk/portal)

<http://www.cs.man.ac.uk/~neville/COMP25111/Lecture11full/Lecture11full.html>



## Where to find this Lecture 11 of the COMP25111 course?

First Go to Blackboard 9; then select: **COMP25111 Operating Systems**

**Week 7**

Then select:

This topic provides...  
10: Memory management 1 (Introduction to basics) by RN;  
11: Memory management 2 (Virtual Memory (1)) by RN;

Then select: **Lecture 11 Information**

Then select: **VIDEOS**



Then select: **Real Time Video of COMP25111 Lecture 11 Tutorial on: Calculation of "Number of Pages" given Address Space and Page size**

**Lecture11Full**

Then select:

© Copyright Richard Neville 2007



1. Question

What is the difference between a 'partition' and a 'program.'

ANSWER(S):

Answer(s):

NOTE: In the exam approximately 2 question are taken from the topics (and program examples) covered in each lecture



# Lecture 11

## Virtual Memory (1)

"The equal-sized blocks concept"

### Paged Virtual Memory

Virtual memory – *"makes the machine appear to have more memory than it actually has."*


Normally this is not the case: virtual = 500G (secondary [hard disk]) & physical = 4G (primary [RAM]).

### Quotation

Pages – *"paging gives the programmer the illusion of large, continuous, linear main memory, the same size as the virtual address space," [1].*

### References:

[1] Quote from: Structured Computer Organization, 5/E, Andrew S. Tanenbaum, Vrije University, Amsterdam, The Netherlands, ISBN-10: 0131485210, ISBN-13: 9780131485211, Publisher: Prentice.



Be able to

- 1) Explain what is meant by a paged virtual memory system;
- 2) Determine the structure of an address in a paged virtual memory system; &
- 3) Establish the outcome of memory references (specified by address) for a paged virtual memory system with a specific page table.
- 4) Self-study NOTES: [Information] No. of bits in Page and Offset (At the end of the lecture notes)

© Copyright Richard Neville 2007

**Footnote**

- 1: Describe: aligned to paraphrase, discuss & give example [pictorial or diagrammatic].  
 2: Determine: aligned to decide upon, find out.  
 3: Establish: aligned to find out, calculate & demonstrate method.

*FOOTNOTE2: D-words: direction words. C-Words: content words.  
 Ref.: Michael J. Wallace (1980, 2004)  
 Study Skills in English, ISBN 9780521537520.  
 D-Words also aligned to; Ref.: Taxonomy of Educational Objectives:  
 The Classification of Educational Goals; pp. 201–207; B. S. Bloom (Ed.)  
 Susan Fauer Company, Inc. 1956.*

Reference: Bloom, B.S., *Taxonomy of Educational Objectives*. Handbook I: The Cognitive Domain. 1956: New York: David McKay Co Inc.

5

**Li** Learning; comprehension; & introspection  
**Virtual Memory**

- Virtual Memory is provided in a computer system for two reasons: Large Add. → Smaller Add.

1) To allow a processor to address a much larger address space than is implemented by the physical memory.

- For example, a modern processor may have a 32-bit address bus, giving access to  $2^{32} = 4\text{GB}$  of memory. In a practical system only 512MB of physical memory may be present.

32 Bits

29 Bits

2) To support the operating system in the management of processes.

- Especially in ensuring that processes do not interfere with one another.

These are the two reasons why Virtual Memory is provided in a computer system.

● Start

● M

● E

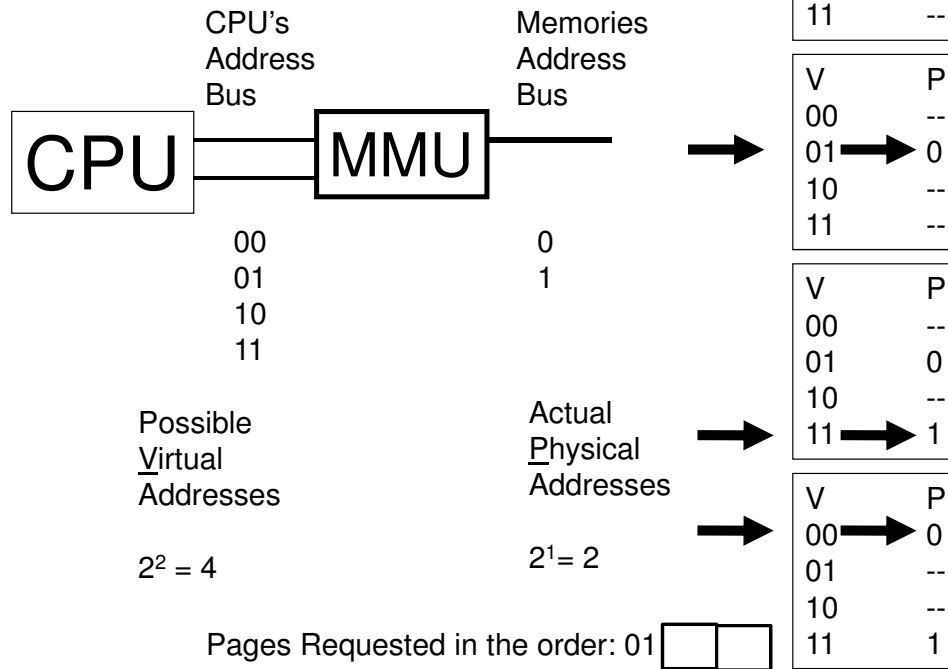
6

Re. **Learning Resources 13.1**, available at the end of the lecture notes.

© Copyright Richard Neville 2007

# Toy Example

© Copyright Richard Neville 2007



7

## Where to find this Lecture11 Virtual And Physical Memory Tutorial one of the Cam Coder Tutorials for the COMP25111 course?

First Go to Blackboard 9; then select: [COMP25111 Operating Systems](#)

### Week 7

Then select:

This topic provides...  
10: Memory management 1 (Introduction to basics) by RN;  
11: Memory management 2 (Virtual Memory (1)) by RN;

Then select:


[Lecture 11 Information](#)

and then: [VIDEOS](#)

Then select:

[Cam Coder Tutorial Lecture11 Virtual And Physical Memory](#)

Then select:



© Copyright Richard Neville 2007

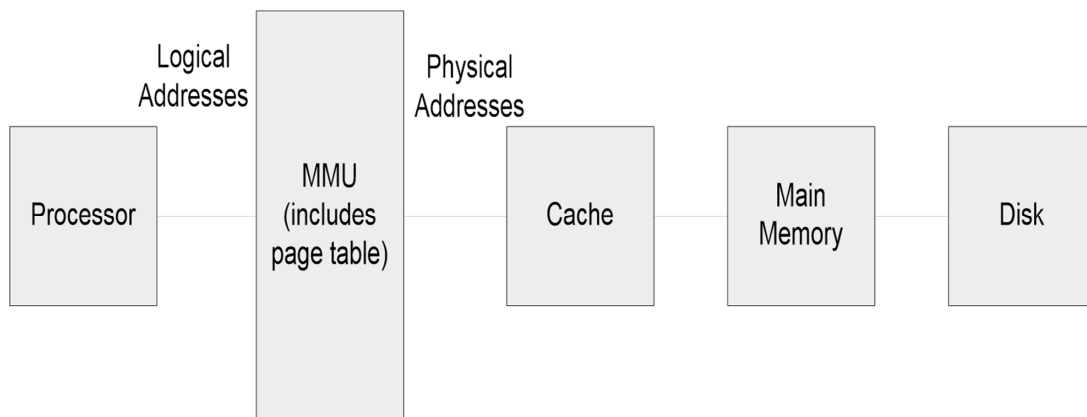
## Memory Management Unit (MMU)

- The MMU is hardware that translates a logical address provided by the processor into a physical address into main memory.
- The MMU is hardware that is placed between the processor and the memory system.
- It uses a table (the page table) to translate from a logical to a physical address.
- The page table is maintained by the operating system.
- The MMU must carry out the translation very quickly because the address of all memory reads and writes are routed through it.

© Copyright Richard Neville 2007

9

## Memory Management Unit (MMU)



© Copyright Richard Neville 2007

10

# Virtual Memory Methods

A few additions to notes are required.

- There are two major methods for implementing virtual memory:

## 1) Paged Virtual Memory

- The pages are of fixed size
  - 512 B to 64 KB are common

2)

- The segments are of **variable** size
  - between 1 and  $2^{16}$  or  $2^{32}$  bytes

# Paged Virtual Memory

- The virtual address space is divided into a number of equal-sized pages (blocks of memory).

**'Page'** in virtual address space

- The virtual address space is given by the number of **bits** in the address bus of a processor:
- So, if a processor has a **32-bit address**, the virtual address space is  **$2^{32}$  Bytes** or **4 GB**.

## Paged Virtual Memory, cont...

- The physical address space is divided into a **smaller** number of page frames, each of which can hold a page.

**'Page frame'** in physical address space

- Thus a page frame is the **same size as a page**.
- In the implementation of a computer system, the physical memory will often be **smaller** than the virtual address space.
- For example, 256 MB may be available when the processor's virtual address space is 4GB.

© Copyright Richard Neville 2007

### Quotation


Page frames – *"pieces of main memory into which the pages go are called page frames."*

References 1: Quote from: Structured Computer Organization, 5/E, Andrew S. Tanenbaum, Vrije University, Amsterdam, The Netherlands, ISBN-10: 0131485210, ISBN-13: 9780131485211, Publisher: Prentice.

13

## Paged Virtual Memory (1)

- QUESTION: How many pages in the virtual address space ?
- Virtual Memory



© Copyright Richard Neville 2007

14



# Paged Virtual Memory (1)

1) The virtual address space will consist of a number of pages:

$$\text{Number of pages} = \frac{\text{Address space}}{\text{Page size}}$$

Given a virtual memory 'address space' of 4GB and a 'page size' of 64KB:

$$= \frac{4 \text{ GB}}{64 \text{ KB}} = \frac{2^{32}}{2^{16}} = 2^{16}$$

e.g. if the virtual address space is 4 GB and the page size is 64 KB there are:

**(64K)** pages.

$$2^{32} = 4,294,967,296 \text{ or } 4\text{G}$$

$$2^{16} = 65,536 \text{ or } 64\text{K}$$

**L** Learning; comprehension; & introspection

## Aside: on units

- When referring to units that are derived from the binary number system, the following are used:


$2^0 = 1$	$2^{10} \sim 1\text{K}$	$2^{20} \sim 1\text{M}$	$2^{30} \sim 1\text{G}$
$2^1 = 2$	$2^{11} \sim 2\text{K}$	$2^{21} \sim 2\text{M}$	$2^{31} \sim 2\text{G}$
$2^2 = 4$	$2^{12} \sim 4\text{K}$	$2^{22} \sim 4\text{M}$	$2^{32} \sim 4\text{G}$
$2^3 = 8$	$2^{13} \sim 8\text{K}$	$2^{23} \sim 8\text{M}$	$2^{33} \sim 8\text{G}$
$2^4 = 16$	$2^{14} \sim 16\text{K}$	$2^{24} \sim 16\text{M}$	$2^{34} \sim 16\text{G}$
$2^5 = 32$	$2^{15} \sim 32\text{K}$	$2^{25} \sim 32\text{M}$	$2^{35} \sim 32\text{G}$
$2^6 = 64$	$2^{16} \sim 64\text{K}$	$2^{26} \sim 64\text{M}$	$2^{36} \sim 64\text{G}$
$2^7 = 128$	$2^{17} \sim 128\text{K}$	$2^{27} \sim 128\text{M}$	$2^{37} \sim 128\text{G}$
$2^8 = 256$	$2^{18} \sim 256\text{K}$	$2^{28} \sim 256\text{M}$	$2^{38} \sim 256\text{G}$
$2^9 = 512$	$2^{19} \sim 512\text{K}$	$2^{29} \sim 512\text{M}$	$2^{39} \sim 512\text{G}$

K = Kilo      M = Mega      G = Giga

0	1.00
1	2.00
2	4.00
3	8.00
4	16.00
5	32.00
6	64.00
7	128.00
8	256.00
9	512.00
10	1,024.00
11	2,048.00
12	4,096.00
13	8,192.00
14	16,384.00
15	32,768.00
16	65,536.00
17	131,072.00
18	262,144.00
19	524,288.00
20	1,048,576.00
21	2,097,152.00
22	4,194,304.00
23	8,388,608.00
24	16,777,216.00
25	33,554,432.00
26	67,108,864.00
27	134,217,728.00
28	268,435,456.00
29	536,870,912.00
30	1,073,741,824.00
31	2,147,483,648.00
32	4,294,967,296.00
33	8,589,934,592.00
34	17,179,869,184.00
35	34,359,738,368.00
36	68,719,476,736.00
37	137,438,953,472.00
38	274,877,906,944.00
39	549,755,813,888.00

## Paged Virtual Memory (1)

- QUESTION: How many page frames in the physical address space ?
- Physical Memory



© Copyright Richard Neville 2007

17

## Paged Virtual Memory (2)

- 2) The physical address space will consist of a smaller number of page frames:

$$\text{Number of page frames} = \frac{\text{Address space}}{\text{Block size}}$$

Given a physical memory 'address space' of 256 MB and a 'block size' of 64KB:

$$= \frac{256 \text{ MB}}{64 \text{ KB}} = \frac{2^{28}}{2^{16}} = 2^{12}$$

e.g. if the physical address space is 256 MB and the block size is 64 KB there are:

**(4K)** page frames.


$$\begin{aligned} 2^{28} &= 268,435,456 \text{ or } 256 \text{ M} \\ 2^{16} &= 65,536 \text{ or } 64 \text{ K} \\ 2^{12} &= 4096 \text{ or } 4 \text{ K} \end{aligned}$$

© Copyright Richard Neville 2007

18

# Paged Virtual Memory

- The logical address generated by the processor is split into two bit fields:



- The number of bits in the page number are given by:
- $\log_2(nv)$
- where
- $nv$  = number of pages in the virtual address space.
- So, for our example of 4 GB of virtual address space and a page size of 64 KB, there are  $2^{16}$  pages (previous slide) giving  $\log_2(2^{16}) = 16$  bits

© Copyright Richard Neville 2007

19

$$2^{16} = 65,536 \text{ or } 64K$$

# Paged Virtual Memory

- The number of bits in the offset is:

- $\log_2(bs)$
- Where  $bs$  = block size.

$$2^{16} = 65,536 \text{ or } 64K$$

- So in our example, with a 64 KB page size there are  $\log_2(2^{16}) = 16$  bits in the offset.
- So address is treated as:



© Copyright Richard Neville 2007

20

# Paged Virtual Memory Procedure

- This procedure can be viewed as a sequence of steps:



- The processor generates a logical address.
- The page number field is used by the MMU to look to see whether the page is in memory or not.
- If it is in memory, a physical address is computed by replacing the page number with the page frame number of where the page can be found.

Together with the offset this is used as a physical address to memory.

- If it is not in memory, the transfer is aborted (page fault) and the operating system will load the page from disk to memory.

## Where to find this Lecture11 Page Table Tutorial one of the Cam Coder Tutorials for the COMP25111 course?

First Go to Blackboard 9; then select: [COMP25111 Operating Systems](#)

### Week 7

Then select:

This topic provides...  
10: Memory management 1 (Introduction to basics) by RN;  
11: Memory management 2 (Virtual Memory (1)) by RN;

Then select:


[Lecture 11 Information](#)

and then: [VIDEOS](#)

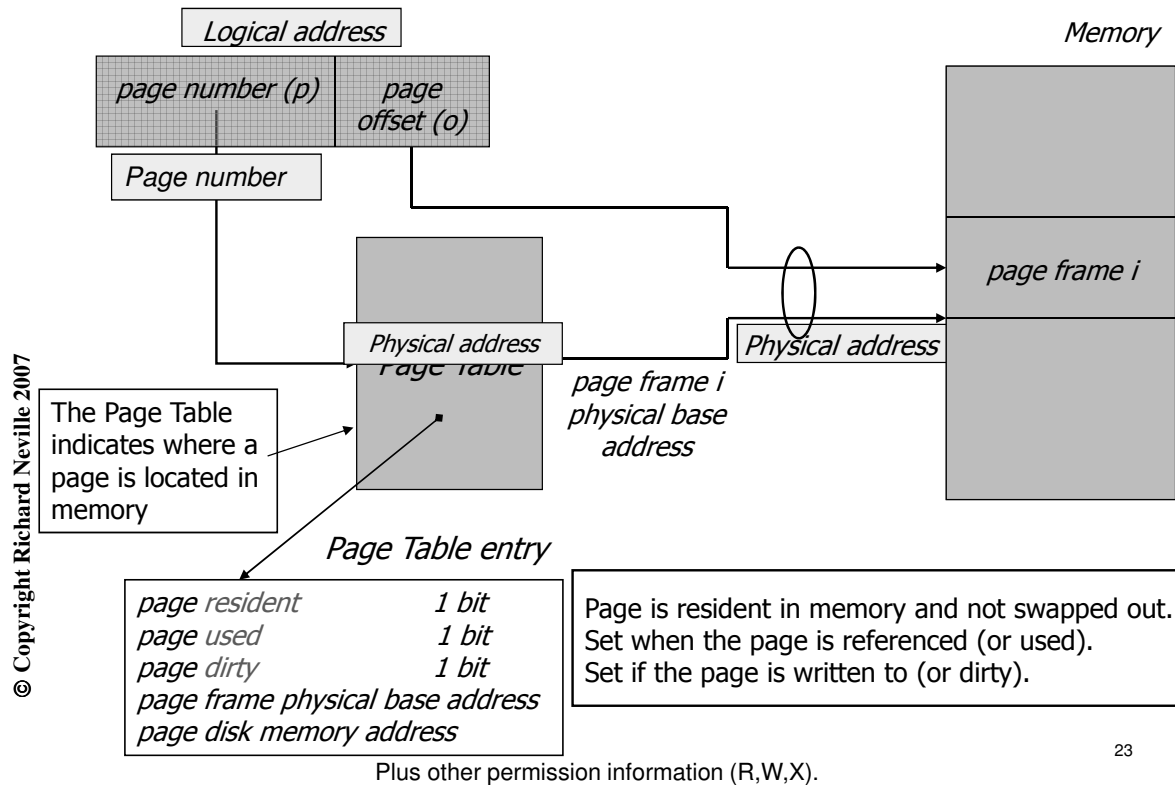
Then select:

[Cam Coder Tutorial Lecture11 Page Table Tutorial](#)

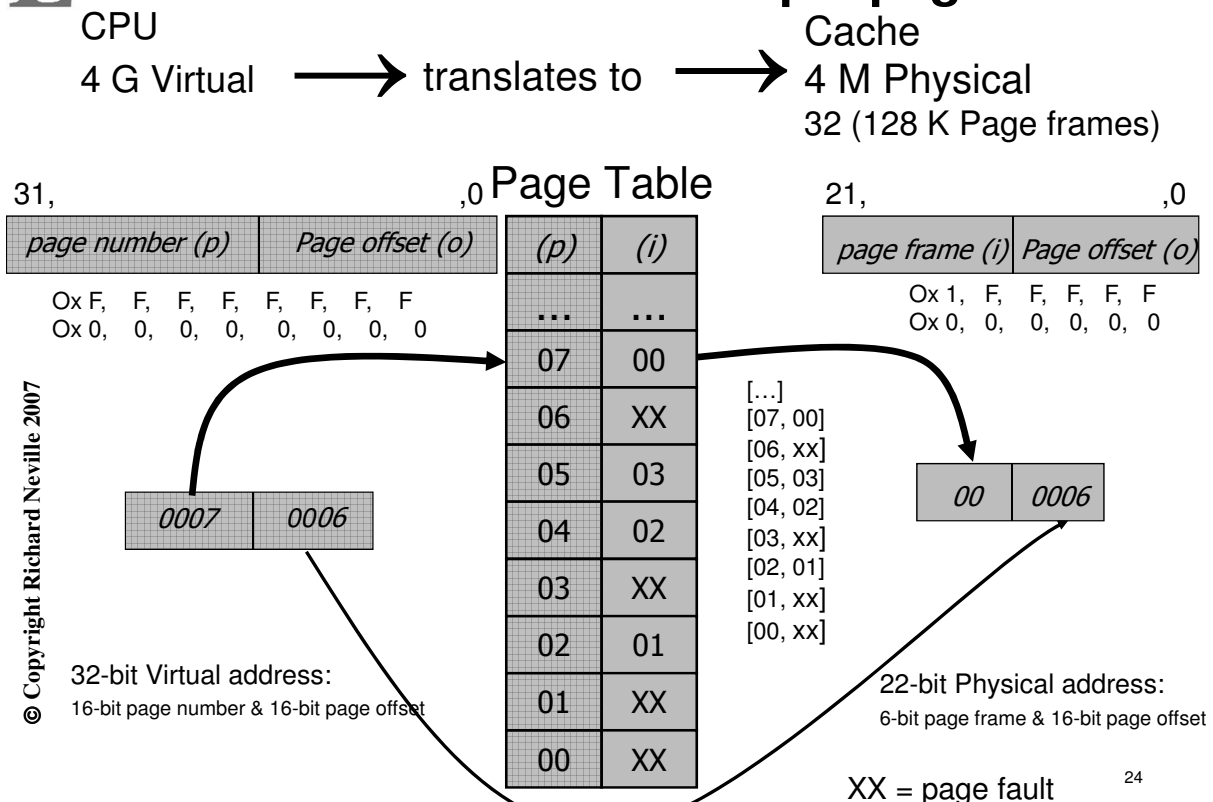
Then select:



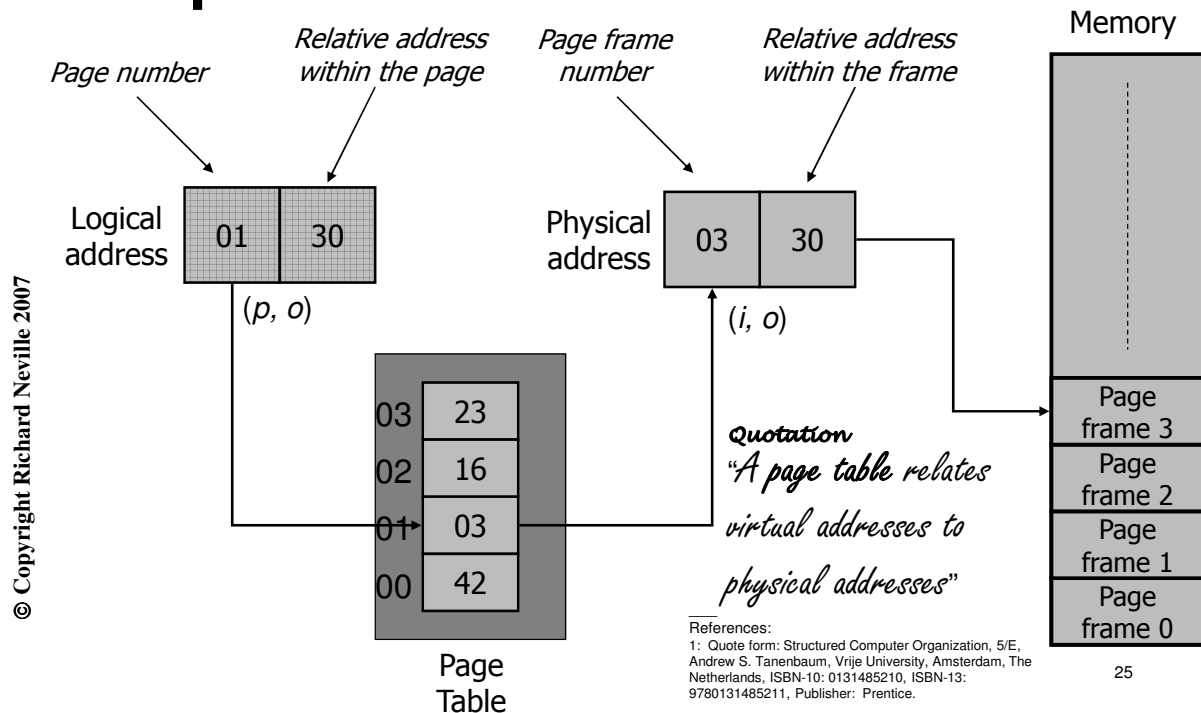
# Paged Virtual Memory Address Mapping



## Simple page table



# Logical to Physical Address Simple View



## Demand Paging

A few additions to notes are required.

- When a reference is made to a page that is not currently in the memory, a page fault occurs.
- This is like an interrupt; initiating the following:
  - 1) The current process will be halted;
  - 2) The operating system will cause the page to be loaded from disk into a page frame;
    - Whilst waiting for the disk, other processes can run in a multitasking operating system.
  - 3) When the page has been loaded, the memory reference will be tried again.

## Page replacement

© Copyright Richard Neville 2007

- After the computer has been operating for a while, each of the page frames in memory will be occupied by a page.
- When another page needs to be loaded, one of the pages already loaded will need to be overwritten.
- An algorithm will need to be used to determine which page is to be replaced:
  - LRU and FIFO algorithms are commonly used.

27

## Page Change

© Copyright Richard Neville 2007

- When the processor writes to memory, the contents of the page will change.
- We need to decide when this needs to be written back to the disk;
  - This is similar to the write strategies for caches.
- In the case of virtual memory, it makes no sense to use a write-through strategy because the time taken to write to disk is so high, so a write-back strategy is used with pages only being written to disk if they have changed and when they are unloaded
- Each page table entry has a dirty bit that indicates whether the page has been written to or not:
  - If it has been written to, the page on the disk needs to be updated when the page is unloaded from the physical memory;
  - If a page has not been written to, there is no need to update the page's data held on the disk.

28

# Simple Example of Page Table

Virtual page number	Page resident	Page dirty	Page frame number
0	1	0	4
1	1	1	7
2	0	0	-
3	1	0	2
4	0	0	-
5	1	1	0


© Copyright Richard Neville 2007

29



# Summary

*Paged VIRTUAL MEMORY*




© Copyright Richard Neville 2007


30

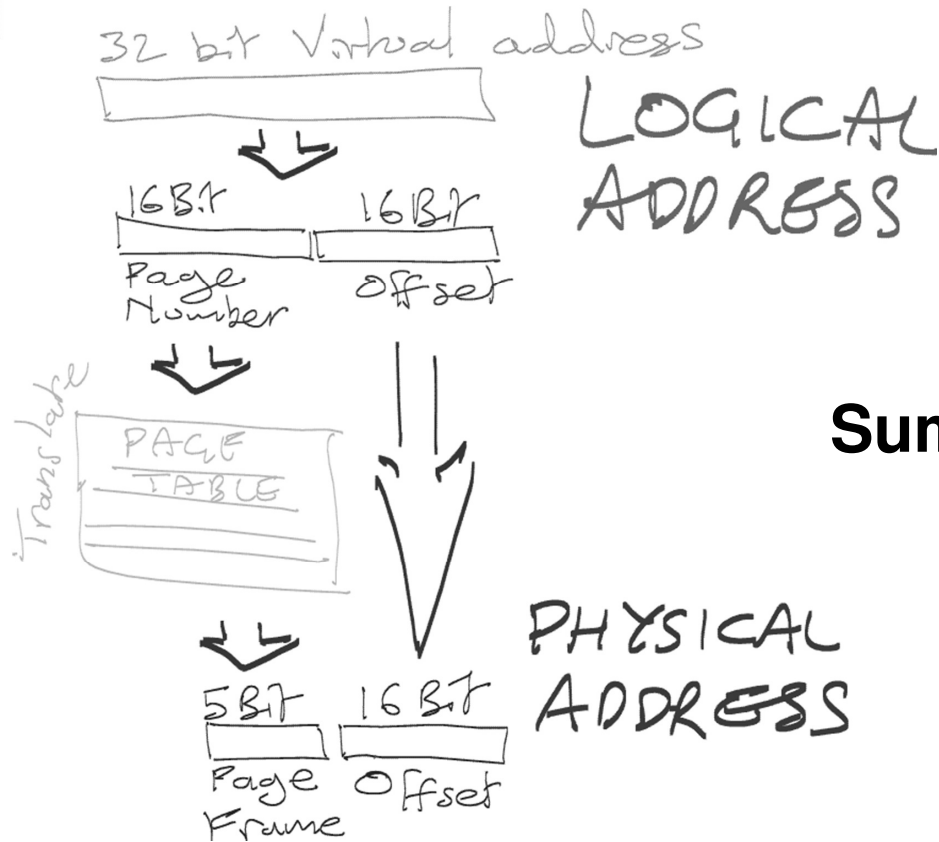


# Summary



# Summary



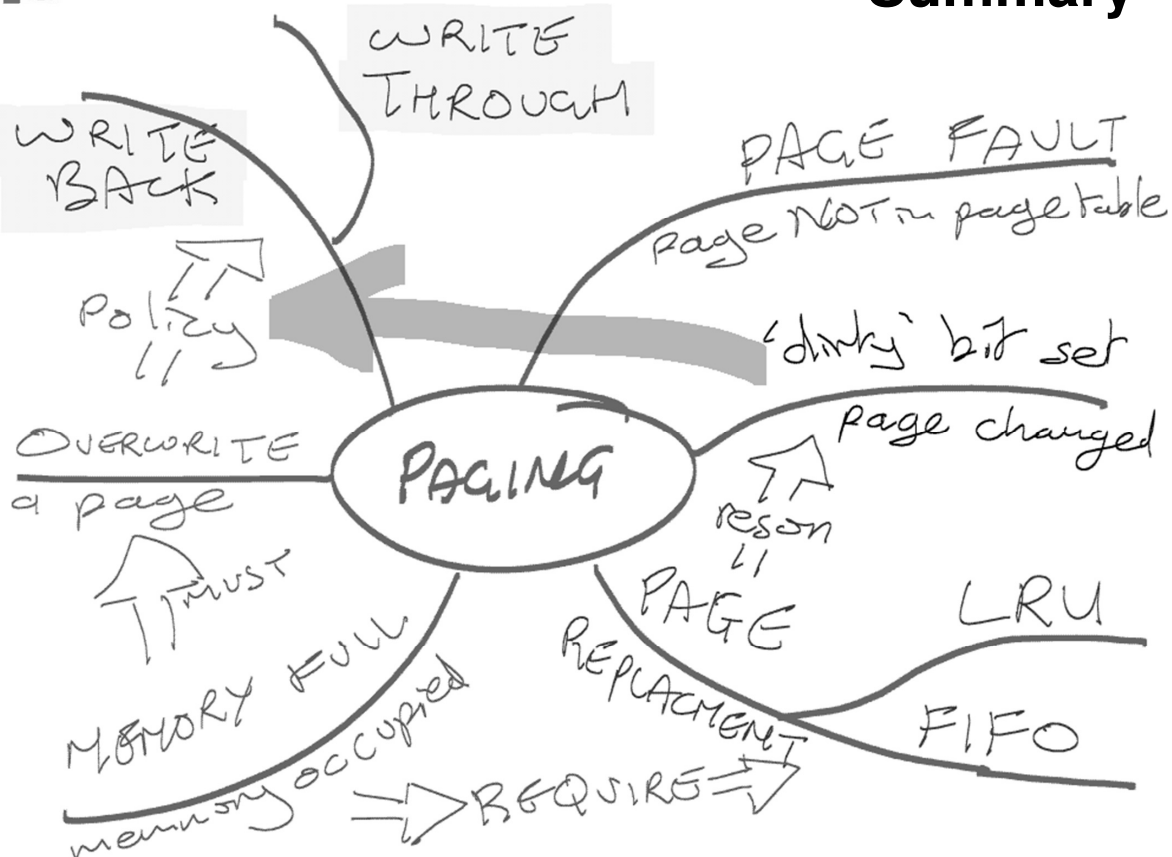


## Summary

33



## Summary





## List of Questions to ask lecturer

- Before the 9a.m. start lecture the lecturer will be half an hour early and you can ask [any and all] questions in that half hour; before the lecture:

- 1.
- 2.
- 3.
- 4.
- 5.

This material is presented to ensure timely dissemination of Lecture materials. Copyright and all rights therein are retained by authors (i.e. © Copyright Richard Stuart Neville.). All persons copying this information are expected to adhere to the terms and constraints invoked by each author's copyright. These works may not be reposted without the explicit permission of the copyright holder © Copyright Richard Stuart Neville.

Personal use of this material is permitted. However, permission to reprint/republish this material for advertising or promotional purposes or for creating new collective works for resale or redistribution to servers or lists, or to reuse any copyrighted component of this work in other works must be obtained from © Copyright Richard Stuart Neville.

Proper referencing of this material is essential. The expected norm for referencing the material is:

*Citation [reference in body of text] (Harvard style):* (Neville, 2010).

*Reference:* Neville, R., (2010). Lecture notes (and all associated materials) for COMP17022 Introduction to Computer Systems; Lecture series, developed and presented by R. Neville.



## Getting ready for next week Do next week's Q3's NOW

- Once you have re-read the lecture notes; and listened to the audio recording [while stepping through the PPT] of the lecture again:
- Please have a think about next week's Q3's
  - on the next page
- If you try to answer the Q3's now you will be in a much better position to recall the information.
- Once you have done this, transfer your answers to next weeks "Student [OWN answers] version" at the start of next weeks lecture.
  - YES this implies bringing the last weeks lecture notes to the next lecture ...



1. Question

What is the difference between a '*partition*' and a '*program*.'

Answer(s):

2. Question

What is a '*fixed partition*'?

Answer(s):

3. Question

Differentiate between fragmentation and compaction:

a) Fragmentation: xxx. B) Compaction: xxx.

Answer(s):

37

NOTE: In the exam approximately 2 question are taken from the topics (and program examples) covered in each lecture



1. Question

Differentiate between the write-through strategy and the write-back strategy.

Answer(s):

2. Question

What does the dirty bit indicate. State how it is utilised.

Answer(s):

3. Question

Discuss the differences between '*page fault*' and '*page demand*.'


Answer(s):

NOTE: In the exam approximately 2 question are taken from the topics (and program examples) covered in each lecture

38

# Self-study NOTES: [Information]

## No. of bits in Page and Offset



**'Page number'** No. of bits =  $\log_2(nv)$ , **'Offset'** No. of bits =  $\log_2(bs)$

© Copyright Richard Neville 2007

Page Number		TOTAL	Offset	
No of Pages ( <i>nv</i> )	Bits in Page No		Bits in Offset	Block size ( <i>bs</i> )
16 pages	4	12 {11,...,0}	8	256 B
256 pages	8	16 {15,...,0}	8	256 B
16 pages	4	20 {19,...,0}	16	64 KB
256 pages	8	24 {23,...,0}	16	64 KB
64 K pages	16	32 {31,...,0}	16	64 KB

39

## GLOSSARY

Using the on-line resources and any other resources compile a glossary of the terms below [PIPLINES: memory management]:

- Virtual memory →
- Paging →
- Virtual address →
- Paged virtual memory →
- Logical address →
- Physical address →
- MMU →
- Page table →
- Translation →
- Segmented Virtual Memory →
- Virtual address space →
- Physical address space →
- Page frames →
- Page number →

© Copyright Richard Neville 2007

40



# GLOSSARY

- [Page] offset →
- Page fault →
- Page resident →
- Page used →
- Page dirty →
- Page frame number →
- Relative address →
- Page table →
- Demand paging →
- Overwritten →
- LRU →
- FIFO →
- Write-through →
- Write-back →
- Dirty bit →



# GLOSSARY & AUX. DATA

- Longer pipeline: ( effect on the processor's performance? & effect of control hazards).
  - In theory the performance improves by  $n$  where  $n$  is pipeline length.
  - However, **control hazards** will have a greater effect in a deeper pipeline, not all of which will be mitigated by branch prediction.  
*There are a number of other comments that could be made about the pipeline in terms of clock speed with longer stages, the ability to fragment instruction execution into so many stages, etc.*
- How does a pipelined processor works: (state what performance improvements result).
  - The technique of processing multiple parts of an instruction at the same time.
  - A sequence of instruction units [stages] which performs a task in several steps, like an assembly line in a factory. Performance increases as the units operate in parallel and this increases the process speed.

# Learning Resources 1

- **Descriptions [Theory] (in text books)**
- Remember the key issues, highlighted in GREEN, are the concepts to look for in any book:
  - Section on virtual memory, paging, virtual address space, MMU, translation, page frame, page fault, demand paging, LRU, FIFO, write-back, & write-through in chapter 9 the operating systems in: Chalk BS, Carter AT, Hind RW (2004) Computer Organisation and Architecture: An introduction 2nd Edition, Palgrave, ISBN 1-4039-0164-3 .
  - Section on virtual memory, paging, memory management, translation, page fault, demand paging – in chapter 7 the operating system support in: Computer Organization and Architecture, Fifth Edition by William Stallings.
  - Section on virtual memory, paging, virtual address space, MMU, translation, page frame, page fault, demand paging, LRU, FIFO, write-back, & write-through – in chapter 6 the operating system machine level in: Structured Computer Organization, 5/E, Andrew S. Tanenbaum, Vrije University, Amsterdam, The Netherlands, ISBN-10: 0131485210, ISBN-13: 9780131485211, Publisher: Prentice. <sup>43</sup>

# Learning Resources 2

- **Descriptions [Theory] (on the web)**
- **Web resources:**
  - Virtual Memory ; available [on-line] @ <http://courses.cs.vt.edu/csonline/OS/Lessons/VirtualMemory/index.html>
  - **Virtual memory simulator** available [on-line] @ <http://www.ecs.umass.edu/ece/koren/architecture/Vmemory/try.html>
  - MOS Free e-book [Low resolution (Not high quality graphics or printing – but readable)]: Modern Operating Systems (MOS) 2nd Edition Andrew Tanenbaum, Available [on-line] @: [http://www.freebookzone.com/fetch.php?bkcls=os\\_thry&bkidx=35](http://www.freebookzone.com/fetch.php?bkcls=os_thry&bkidx=35)
  - Virtual Memory II, part of course: COMP3231/9201/3891/9283 Operating Systems 2011/S1, Available [on-line] @: <http://cgi.cse.unsw.edu.au/~cs3231/11s1/lectures/lect15.pdf>



# Questions

## Introduction to Questions:

The set of questions are based on lecture 13.

Answer Sheet will be given later in year and will contain the answers to these questions.

© Copyright Richard Neville 2007

- Remember to find detailed and comprehensive answer you should [also] reference associated text books in the library.
- A reasonable starting place for associated book titles are:
  - 1) This units 'module guide'; given to you in RN's first lecture – or on the web [Blackboard];
  - 2) Those books mentioned in 'Background Reading;'
  - 3) Those books [and web resources] mentioned in Learning Resources.

45



# Questions

## 1. Question:

- Given the following address spaces and associated page sizes; calculate the number of pages that result in the *virtual address space*:
  - a) Page size 32KB and virtual address space 2GB;
  - b) Virtual address space 8GB and page size 256KB.

Answer:

© Copyright Richard Neville 2007

46





## 1. Answer:

- Given the following address spaces and associated page sizes; calculate the number of pages that result in the *virtual address space*:
  - Page size 32KB and virtual address space 2GB;

Answer(s):



## 2. Question:

Given the set of specific address size and associated block size below. Calculate the number of page frames in the physical address space :

- Physical address space 1 GB and the block [page] size of 64 KB; and
- Given a block [page] size of 128 KB and a physical address space of 2 GB.

Answer:



## 2. Answer:

Given the set of specific address size and associated block size below.  
Calculate the number of page frames in the physical address space :

- a) Physical address space 1 GB and the block [page] size of 64 KB;

Answer(s):



## 3. Question:

Given the page table, on the right, and the set of virtual addresses  $(p, o)$  below. Answer the following questions for each of the virtual address:

- Will the virtual address generate a page fault?
- If it does, generate a page fault, what must the MMU do? [If not; what does it do then?]
- What is the physical address; if 'no' to a)?

3.1)  $(p, o) = \text{Ox } 00020006$ ; &

3.2)  $(p, o) = \text{Ox } 0006000F$ .

Given  $p$  is a 16-bit number and  $o$  is a 16-bit number.

## Answer:

Page table

$(p)$	$(i)$
...	...
07	00
06	XX
05	03
04	02
03	XX
02	01
01	XX
00	XX

Note1 : the page table only displays the lower 2 Hex digits, of the page number, not all four.

Note 2: XX denotes no page frame number allocated, yet!



## 3. Answer:

Given the page table, on the right, and the set of virtual addresses (p, o) below. Answer the following questions for each virtual address:

3.1) (p, o) = 0x 00020006

a) Answer(s):

b)

c)

•

3.2) (p, o) = 0x 0006000F

a) Answer(s):

•

b)

•

c)

•

Page table

(p)	(i)
...	...
07	00
06	XX
05	03
04	02
03	XX
02	01
01	XX
00	XX

51



## 4. Question: exercise in conversion from page to page frame number

- A computer system operates using a paged virtual memory system (no segmentation). The processor address bus is 40 bits wide, the main memory in the system is 1GB and the page size is 128KB.

- What is the size of the virtual address space?
- How many pages are there in the virtual memory?
- How many page frames are there in the main memory?
- The 40-bit virtual address will be split into two bit fields, the *page number* and the *offset*, state how many bits there will be in each field.
- Suppose that part of the page table is given below (entries for 'page frame' number given in hexadecimal representation of a 23-bit binary value). Where it is possible to determine the physical address, state what this will be for the following addresses (given in hexadecimal):

- 0x000000ACEF
- 0x000000B020
- 0x0000004FCDA

- Now, given the Page table opposite; which of the 'page frame Number(s)' are Viable? If not why?

Virtual Page Number	Page Resident	Page Dirty	Page Frame Number
0	1	0	3AAC00
1	1	1	0456BFC
2	0	0	-
3	1	0	000200
4	0	0	-
5	1	1	01576F

52



## 4. Answer

© Copyright Richard Neville 2007

a.

Answer(s):

b.

c.

53



## 4. Answer

© Copyright Richard Neville 2007

d. There are  $2^{23}$  pages to be addresses, this we

Answer(s):

54



## 4. Answer e.

# Questions

i. Firstly split the address into the frame number and offset:

1) 000000ACEF hex

2) Answer(s):

3)

4)

5)

6)

7)

8)



## 4. Answer e.

# Questions

ii. Firstly split the address into the frame number and offset

1) 00000B0020 hex

2) Answer(s):

3)

4)

5)

6)

7)

8)



## 4. Answer e.

iii. Firstly split the address into the frame number and offset

- 000004FCDA hex

• Answer(s):

•

•

•



## 4. Answer f.

Answer(s):



## 5. Answer:

Explain the difference between a virtual memory address and a physical memory address.

Answer(s):



6. A 16 bit [virtual] address is divided into an 8-bit page number and an 8-bit page offset; whereas the associated physical address a 7-bit page frame with the appropriate page offset address sizing. Show, with the aid of a diagram, the table structure necessary to convert this virtual address into a real [physical] address.

**Answer:**



Learning; comprehension; & introspection

Long [& Short] Exam Questions

## Questions

6. A 16 bit [virtual] address is divided into an 8-bit page number and an 8-bit page offset; whereas the associated physical address a 5-bit page frame with the appropriate page offset address sizing. Show, with the aid of a diagram, the table structure necessary to convert this virtual address into a real [physical] address.

**Answer:**

Answer(s):



Learning; comprehension; & introspection

Long [& Short] Exam Questions

## Questions

7. What is a 'page fault'? Describe how a page fault is handled by the Memory Management Unit and the operating system..

**Answer:**

Answer(s):





## Revision Exercises

- Scan read Lecture 13's Questions.
  - Answer Lecture 13's Questions
    - Particularly those questions you had difficulties with when you first tried them.



## Background Reading

- [1] Chapter 9 the operating systems in: Chalk BS, Carter AT, Hind RW (2004) Computer Organisation and Architecture: An introduction 2nd Edition, Palgrave, ISBN 1-4039-0164-3 .
- [2] Chapter 7 the operating system support in: Computer Organization and Architecture, Fifth Edition by William Stallings.
- [3] Chapter 6 the operating system machine level in: Structured Computer Organization, 5/E, Andrew S. Tanenbaum, Vrije University, Amsterdam, The Netherlands, ISBN-10: 0131485210, ISBN-13: 9780131485211, Publisher: Prentice.
- [4] Operating system concepts, A. Silberschatz, P.B. Galvin, & G. Gagne, Wiley;
  - Chapter 9 virtual-memory management.
- [5] Modern Operating Systems, Andrew S. Tannenbaum, Various Editions, Prentice Hall;
  - Chapter 4 Memory management.

## COMP25111 Exercise 3:

### Simulation of Paging Behaviour

Contains a copy of the “traceA” file.

Before you start this PLEASE read the multimedia PPTs and audios on the Blackboard 9 portal- on the Blackboard 9 Web Site; select: COMP25111 Operating Systems  
<https://www.portal.manchester.ac.uk/>; then  
1/ Go to COMP25111 Operating Systems 2011-12 1st Semester;  
2/ Then Lab3\_RN  
3/ Then Lab3 Hints  
4/ Then read through as many hints as you need; remember the methods, processes, naming conventions, and implicit theory outlined in the hints will be useful to you in your final year projects where you should think about using these techniques...

**Duration: 1 Session**

#### 1. Learning Outcomes

On completion of this exercise, a student will:

- Have implemented the functionality of a simple Memory Management Unit (MMU) for a paged virtual memory system in Java.
- Have exercised the MMU using a trace file of memory accesses using a Java program which initialises the MMU and calls methods within it to perform read and write actions.
- Have produced statistics, which show the behaviour of the MMU with particular reference to page faults and page rejections.

#### 2. Introduction

Virtual memory is a technique for providing a process with apparent access to the whole addressable memory space of a processor in circumstances where the real memory available may be considerably smaller. One advantage of the scheme is the process can use larger code and data sizes than would otherwise be possible. Another is the run-time layout of code and data can be considerably simplified if it does not have to fit into a confined space. This is particularly true for dynamic data whose size is unknown until run-time.

Virtual memory relies on the fact that, in the majority of programs, the use of both code and data exhibits both *spatial locality* and *temporal locality*. That is that at any point in time, the program is usually only using a small fraction of its code and data and that, over a short time period, the usage will not change significantly.

A virtual memory system therefore tries to keep currently used code and data in the (limited size) real memory of the system while the rest (if it exists) is kept on background storage such as magnetic disk. In order to relieve the programmer from the complex task of working out which data should be where at any point in time, the system arranges to move the data between real memory and disk automatically as required.

A piece of hardware known as a Memory Management Unit (MMU), together with system software within the Operating System provide all the functionality needed to

do this. The purpose of this exercise is to implement the functionality of a MMU in order to gain an understanding of the principles.

#### 3. Paged Virtual Memory

Devices such as disks usually store and retrieve data in blocks of hundreds or thousands of bytes. Accessing a block is often slow compared to CPU memory speeds but once accessed; the contents of that block can be read or written very rapidly. In these circumstances, it makes sense to move data in units which are larger than bytes or CPU words. Virtual memory systems therefore operate at the unit of a *page* when transferring data to and from memory and disk. A page does not necessarily correspond to a disk block size but is fixed for a particular MMU typically between 4 and 64 kilobytes. This will usually correspond to several disk blocks.




Figure 1 shows the basic structure of a paged virtual memory system.

The addresses issued by the CPU are virtual addresses which are the full width of the processor's addressing capability. In a modern 32 bit processor (for example ARM or x86) these will be a full 32 bits capable of addressing 4Gbytes of memory. The real memory will usually be smaller, typically 512 Mbytes on a laptop.

Each of these memory spaces is considered to be divided into pages. Note that this is not reflected in the internal structure of any memory, it is just an abstract division; the memory addresses are still binary values which cover a linear address space (usually at the unit of bytes). However, we can view any address as being divided into two sections, one which addresses the page and one which is an offset within the page which addresses individual units (bytes). By convention, we usually refer to the virtual address as having *virtual page numbers* (or just page numbers) and the real address as having *page frame numbers*. Figure 2 shows this pictorially.




Figure 2 shows this pictorially view of *virtual page numbers* and *page frame numbers*.

The function of the MMU is to keep track of which virtual pages exist in real memory and, which on disk. It does this by keeping a *page table* of every possible virtual page number which contains the page frame numbers of any real copies. It is important to note that the offset is common to both virtual and real addresses. In order to find the data associated with a virtual address it is only necessary to take the virtual page number and replace it by the page frame number (if it exists). The page frame number is looked up in the page table to produce the real address which can then be used to access real memory. This is usually referred to as *address translation*.

If the table indicates that a virtual page does not exist in real memory this is called a *page fault*. It is necessary for the MMU to initiate a transfer of the page's data from the background storage (disk). The exact mechanism for this is not relevant here but it should be noted that this is a relatively slow process and the CPU itself can organise most of the transfer.




Figure 3 shows more detail of the MMU table and the address translation operation.

The page table is shown here as a single linear structure with the lookup done by the full page number. For large virtual address spaces, this may not be practical and more complex techniques may be required (this is covered in the course lectures).

#### 4. Page Rejection

If there is a page fault, a place needs to be found in real memory for the page copy. Initially, assuming all real memory is not in use, it is possible to keep allocating page frames linearly through the address space. However, if all real memory is in use, we must find a page to reject. We must then write its contents back to disk before reallocating its page frame number to the page access, which has caused the fault.

We need to decide which page to choose using a *page rejection algorithm*. The most commonly used is an algorithm called LRU (least recently used). It works on the simple principle that the page, which has survived the longest time without being accessed, is a good candidate for replacement.

One useful optimisation is to note if a page has been written to since it was bought from disk into real memory. If not, there is no need to write it back to disk before reusing the real page frame. This is often true for pages which contain code.

#### 5. The Exercise

The exercise is composed of two parts:

**Part One:** implements the majority of the code and the LRU page replacement algorithm. The majority of the (theoretical) advice below will aid you in this task;

after you have designed, developed, and implemented the program utilise the two trace files: **traceA** and **traceB**; to test your program.

**Part Two:** implements the a second algorithm the FIFO page replacement algorithm. The advice below will not aid you in this task. Nevertheless, it is expected that once you have implemented (and tested) the LRU algorithm you will have gained sufficient knowledge to enable you to undertake the FIFO design, development and testing on your own. After you have designed, developed, and implemented the program utilise the trace file: **traceLRUandFIFO**; to test your program. The FIFO algorithm has been covered in your lecture series and it is covered extensively in your course books; also see the Learning Resources (book & Web resources) as well as the Background Reading (book & Web resources) at the end of Lecture COMP25111 Operating Systems Lectures 14: Virtual Memory (3) for further reference material.

### Introduction to Exercise

The purpose of the exercise is to write Java code to implement a simple MMU, which contains a page table and code to perform the necessary actions when read, or write accesses occur to a virtual memory address. Because the exercise is using a trace of addresses rather than executing a real program, there is no need to perform any real memory accesses. Instead, it is simply necessary to access the page table to determine if there is a page fault and, if so, find a free page frame to use and place it in the table. If, initially, there are unused page frames these can be allocated linearly until there are none free. At that point it will be necessary to find a page to reject and use its page frame for the newly required page.

You should implement a LRU algorithm to determine page rejection. To make this simple, you should have a 'timestamp' in each page table entry, which is updated appropriately, when a page is accessed.

You should also include a 'dirty' indication in the page table implementation, which is set on a write. This indicates if the page needs to be written back to disk if it is rejected.

Note that a practical page table will usually also contain page access information which determines the types of access which are allowed to a particular page (for example read only). There is no need to include such information for the purposes of this exercise.

In appropriate Blackboard subfolder, you will find a program harness which, after creating a MMU object, reads the trace of store accesses and makes calls to read and write methods within that object until the trace file ends. It then prints out some statistics:

```
Total Accesses
Total Instruction Fetches
Total Page Faults
Total Page Rejections (i.e. ignoring startup page faults)
Total Page Writebacks (to disk)
```

The first two are included in the given program; the last three need to be calculated by your MMU implementation. (A skeleton MMU.java file is provided in the appropriate Blackboard subfolder)

The trace file format is a line for each access. The first number indicates the access type (0=read, 1=write, 2=fetch) and the second is a (hex) **24 bit memory address**. Within the program, the address is considered to be a 12-bit page number and a 12 bit offset representing 4k pages each of 4k bytes, i.e. 16Mbytes in total of virtual memory

space. This is deliberately small so that your implementation can use a simple single level page table. Although obviously not representative of modern systems, this is typical of a virtual memory system in computers of the 1960s when virtual memory was invented (in Manchester University on the Atlas computer [1][2]).

An example of 'traceA' file is given below:

```
2 400000
0 600000
2 400001
1 600001
2 400002
0 700000
2 500000
0 700001
```

The first line "2 400000" is decoded as follows:

First number	Second number
2	400000
2=fetch	0100 0000 0000 0000 0000 0000 <sub>HEX</sub>
This is a fetch access	24 bit hexadecimal memory address

The following decode the first number of each of the eight lines in traceA:

First number	Type of access
2	2=fetch
0	0=read
2	2=fetch
1	1=write
2	2=fetch
0	0=read
2	2=fetch
0	0=read

The program given makes a single call to the method which performs the trace simulation using a very very small real memory (2 [real] page frames, i.e. 8k bytes [in MMUSim.java the real memory sizes (in pages) is set to "rmem\_pages = 2;"). This should be run with the file traceA when initially debugging your implementation. This file contains only 8 store accesses which have been devised so that you should be able to predict what statistics your MMU should produce.

When you think this is working, you should modify the program to perform simulations using both 32 and 64 real page frames on traceB. This represents 128k and 256k bytes of real memory, again realistic in the 1960s.

### 6. Results

You should be prepared to demonstrate the results of the runs on traceA and traceB when your exercise is marked. To help you determine if you have got a correct implementation, the total number of page faults for 32 pages and traceB should be 227.

You should also be prepared to describe your code and draw relevant conclusions from the results.

## Assessment

### COMP20051 Exercise 3: Developing Simulation of Paging Behaviour

Demonstrators please fill in the student's full detail [below] before starting the assessment; may be the best way to do this is to get the students to fill this in themselves:

Students full name	Student Number	Email address	Course studied

At the deadline, at the end of the Session, the marks awarded are as given in table 1.

Demonstrators please assign marks for each of the eight questions in boxes  provided.

At the deadline, at the end of Session 5, the marks awarded are as follows:

Mark awarded for:	Exercise 3
1. <b>Show</b> Correct page table [data] structure; plus comment on how you [the student] developed the data structure.	1 <input type="text"/>
<b>Demonstrating</b> an overall [software engineering] correct approach; of the algorithmic and combined data structure: <b>2 – 5</b> .	
2. <b>State</b> how REQUIREMENT development was undertaken.	1 <input type="text"/>
3. <b>State</b> how [if any] ABSTRACT DESIGN was undertaken.	1 <input type="text"/>
4. <b>State</b> how IMPLEMENTATION was undertaken.	1 <input type="text"/>
5. <b>Demonstrating</b> a methodology for TESTING the program; and comment on: <b>interpretation of results</b> .	1 <input type="text"/>
6. <b>Showing</b> the generation of an appropriate OUTPUT trace file: <b>traceA</b> files [displayed in real-time on command [Prompt] screen] (using LRU).	1 <input type="text"/>
7. <b>Showing</b> the generation of an appropriate OUTPUT trace file: <b>traceB</b> files [displayed in real-time on command [Prompt] screen] (using LRU).	2 <input type="text"/>
8. <b>Showing</b> the generation of an appropriate OUTPUT trace file: <b>traceLRUandFIFO</b> files [displayed in real-time on command [Prompt] screen] (using LRU and FIFO).	2 <input type="text"/>
<b>TOTAL mark</b>	<b>10</b> <input type="text"/>

Remember to save these exercises in a directory nominally named .../COMP20051/ex3.

Do this exercise [or save this exercise] in a directory named COMP251111/ex3 directory, save your downloaded template [Skeleton(s)] code:-

MMUSim.java; and  
MMU.java.

as well as the trace files:

traceA;  
traceB; and  
traceLRUandFIFO

... in COMP25111/ex3.

Remember to save your trace files of your solution.

The three files traceA; traceB; and traceLRUandFIFO; as well as the template [Skeleton(s)] MMUSim.java & MMU.java are downloadable from Blackboard in Folder Lab3.

## References

[1] One-Level Storage System, T. Kilburn, D.B.G. Edwards, M.J. Lanigan, F.H. Sumner, IRE Trans. Electronic Computers April 1962

[2] Tom Kilburn (1956), The Atlas, School of Computer Science: information on Tom Kilburn's [computing] effort known as the MUSE (microsecond) computer , available [on-line] @ <http://www.computer50.org/keill/atlas/atlas.html>. [Last accessed 12/10/09, 11:09].

## 7. Decoding the COMP20051 Laboratory Exercise 3; and ‘supporting advice’

The previous six pages document is a typical laboratory that you are normally set in the School of Computer Science. The first task is to decode [work out] what explicitly is required. As you have read the document once now reread section 5; this time highlighting or extracting the main requirements of the exercise.

### Hint 1: [#H1]

[#H1] If you require a hint, download them from Blackboard. ‘Hint No 1’ hints at the requirements you should have data mined from reread section 5.

After reading document ‘H1.doc’ to validate you have extracted the relevant requirements you next step is to start the design phase that fulfils the requirements [#R1 to #R5] of the exercise. Then once designed the MMU class can be implemented.

### Hint 2: [#H2]

[#H2] If you require a hint, download them from Blackboard. After you have compiled a list of requirements, you should now think about designing the MMU class. ‘Hint No 2’ suggests an approach you could take.

Reading document ‘H2.doc’ may help you move forward in the design phase or validate your design.

You should now be able to implement, test, and run MMUSim utilising traceA. Once you have reached this stage you can compare your program’s trace output with a typical trace A output [result] available on Blackboard; the file is named ‘traceAres.’

One of the final requirements of the exercise is to modify the program [MMUSim] to perform simulations using both 32 and 64 real page frames on traceB. To do this you should consider altering the number of real memory pages.

Finally, you should consider the best way to describe your code; when asked. In addition, you should think about drawing relevant conclusions from the results.

### The “>>” operator

Final point: what does the operator “>>” actually do?

In the code snippet:

```
int addr = access.addr;
```

```
int vpage_no = addr >> 24 - vmem_bits; // assumes a 24 bit virt addr
given vmem_bits = 12. Hence the actual value of ‘24 - vmem_bits’ is “addr >> 12”
```

In “doSimul( ... )” in class MMUSim.Java. The operator “>>” is utilised. If we look at what happen

Values:  
addr = 4194304  
vpage\_no = 1024

So what has happened?

The java operator “>>”

>> shift bits right with sign extension

It implies [with two objects ‘x’ & ‘y’]:

**x >> y**

Means: Shift Right - Signed

Explicitly means: Shift **x** to the right by **y** bits. Low order bits are lost. Same bit value as sign (0 for positive numbers, 1 for negative) fills in the left bits.

Given the example addr = 4194304  
which is specified for addr an integer (int) which is base (or radix) 10.

Converting 4194304<sub>10</sub>  
Converting to hexadecimal:  
00400000<sub>16</sub>  
or radix 2 it is:  
0000 0000 0100 0000 0000 0000 0000<sub>2</sub>

Hence given **addr >> 12**” After shifting right 12 places this is:

0000 0000 0000 0000 0000 0100 0000 0000<sub>2</sub>  
which is:  
00000400<sub>16</sub>  
or radix 2 it is:  
1024<sub>10</sub>

QED