

Transport Services and Protocols (2)

Andy Carpenter

(Andy.Carpenter@manchester.ac.uk)

Elements these slides come from Kurose and Ross, authors of "Computer Networking: A Top-down Approach", and are copyright Kurose and Ross

Transmission Control Protocol (TCP)


- Provides robust process-to-process delivery service
- Creates a virtual connection between applications
- Issues:
 - identifying connections, establishing connections
 - implementing service model, port mechanism
- Implements finely tuned congestion-control mechanism
- Used when application wishes to avoid:
 - complexities of network, error recovery
- Typically used for user applications; e.g.:
 - email (SMTP), web surfing (HTTP)

TCP: Service Model

- Substantially modifies service model of IP; provides:
 - connection-oriented delivery between applications
 - reliable byte stream
 - all bytes sent are delivered in-order
 - unstructured stream; no break into messages
 - full-duplex connection (two streams)
 - buffered; data not immediately sent and delivered
 - each stream is independently flow-controlled

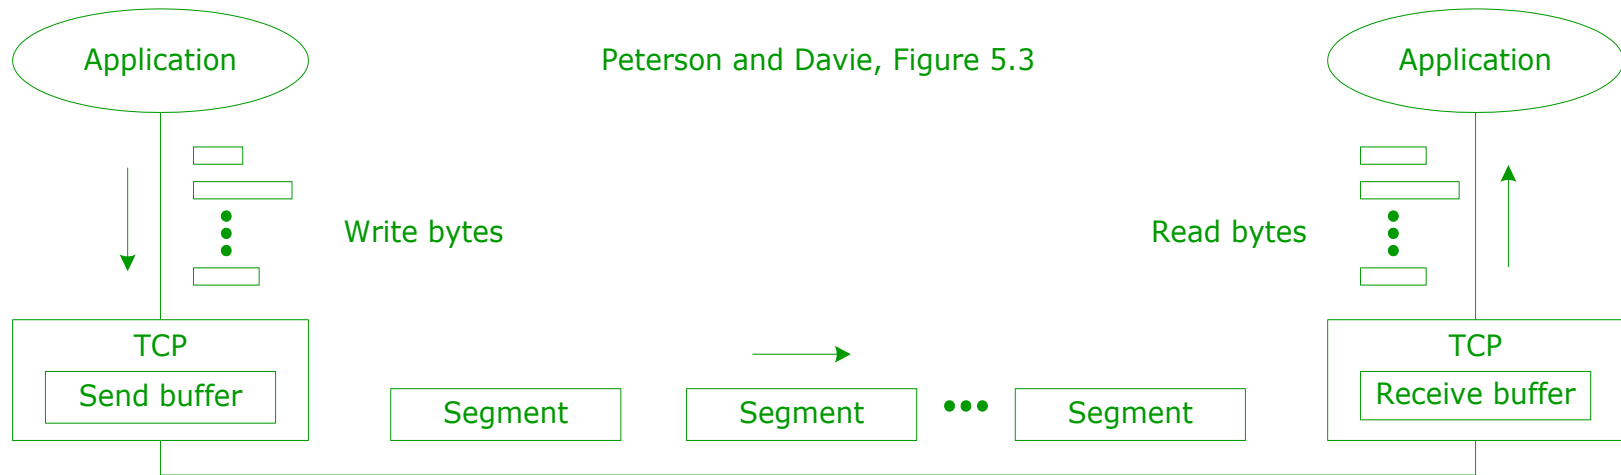
TCP: Identifying Connections

- TCP forms a virtual connection between applications
- Service end of connection may not be unique; e.g.:



- Connection is identified by its end points; value of:
 - source IP address, source port number
 - destination IP address, destination port number

TCP: Segmentation



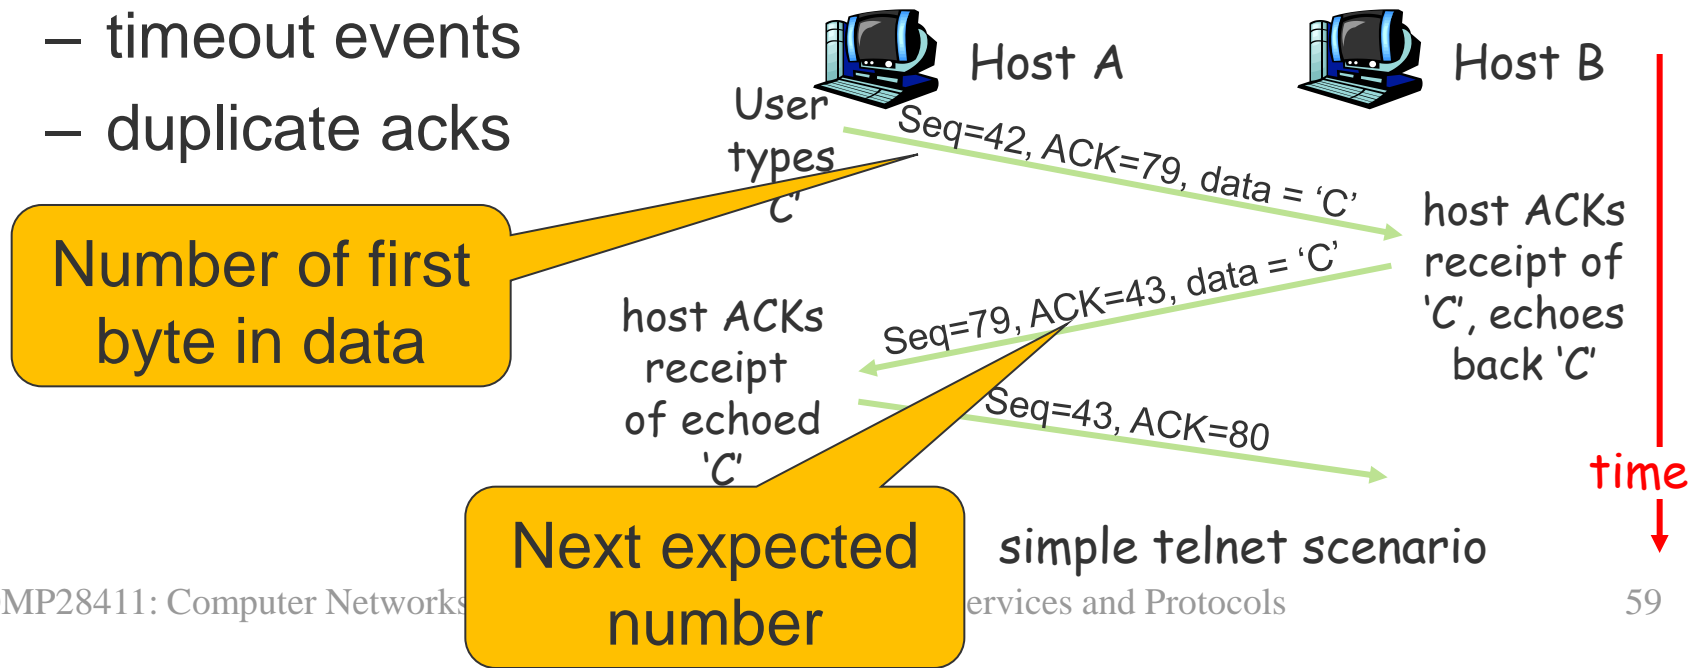
- Application writes bytes to connection as needs
- TCP collects (buffers) bytes before sending
- Sends set of bytes as a TCP segment in IP datagram
- Receiving TCP also buffers
- Application reads bytes from connection as required

TCP: Reliability

- TCP uses
 - sliding window, Go-Back-N (cumulative ACKs)
 - sequence numbers for bytes not segments


Why? single retransmission timer

- Retransmissions are triggered by:
 - timeout events
 - duplicate acks



TCP: Reliability

- Example, window 8, last ACK 7



TCP: Reliability

- Individual bytes have a 32-bit sequence number
- Send sequence number of first byte in segment
- Sequence numbers of other bytes implicitly transmitted
- Sequence numbers in each direction are independent
- Segment numbers must be unique; avoid wraparound
- TCP uses size of identifier and probable data rate
 - at 1MB/sec, about 1 hour to wraparound

TCP: Value of Retransmit Timeout


- If value too small, unnecessarily retransmit segments
- If value too large, get excessive delays before retransmit
- Appropriate value related to RTT, but that depends on:
 - pair of hosts involved, time, congestion in network
- For each connection:
 - use adaptable algorithm to determine RTT
- Basic algorithm sets timeout to twice estimated RTT
- Karn/Partridge algorithm reduces miscalculations
- Jacobson/Karels algorithm
 - copes with significant variance in real RTT

TCP: Estimation of Connection RTT

- Estimate RTT as average of measured RTTs:
- When re-sending, to which transmit does ACK relate?
- Karn/Partridge algorithm:
 - only measures RTT for non-retransmitted segments
 - doubles timeout value for each retransmission
- Jacobson/Karels Algorithm:
 - if RTT variation small, average is good approx.
 - if RTT variation large, average is poor approx.
 - algorithm takes account of variation

TCP: Connection Flow Control

- Receiving buffer has finite capacity; must not overrun



Flow control

Sender will not overflow receiver's buffer by transmitting too much, too fast


Speed-matching service

Matching the send rate to the receiving applications drain rate

Flow control allows receiver to influence transmitter

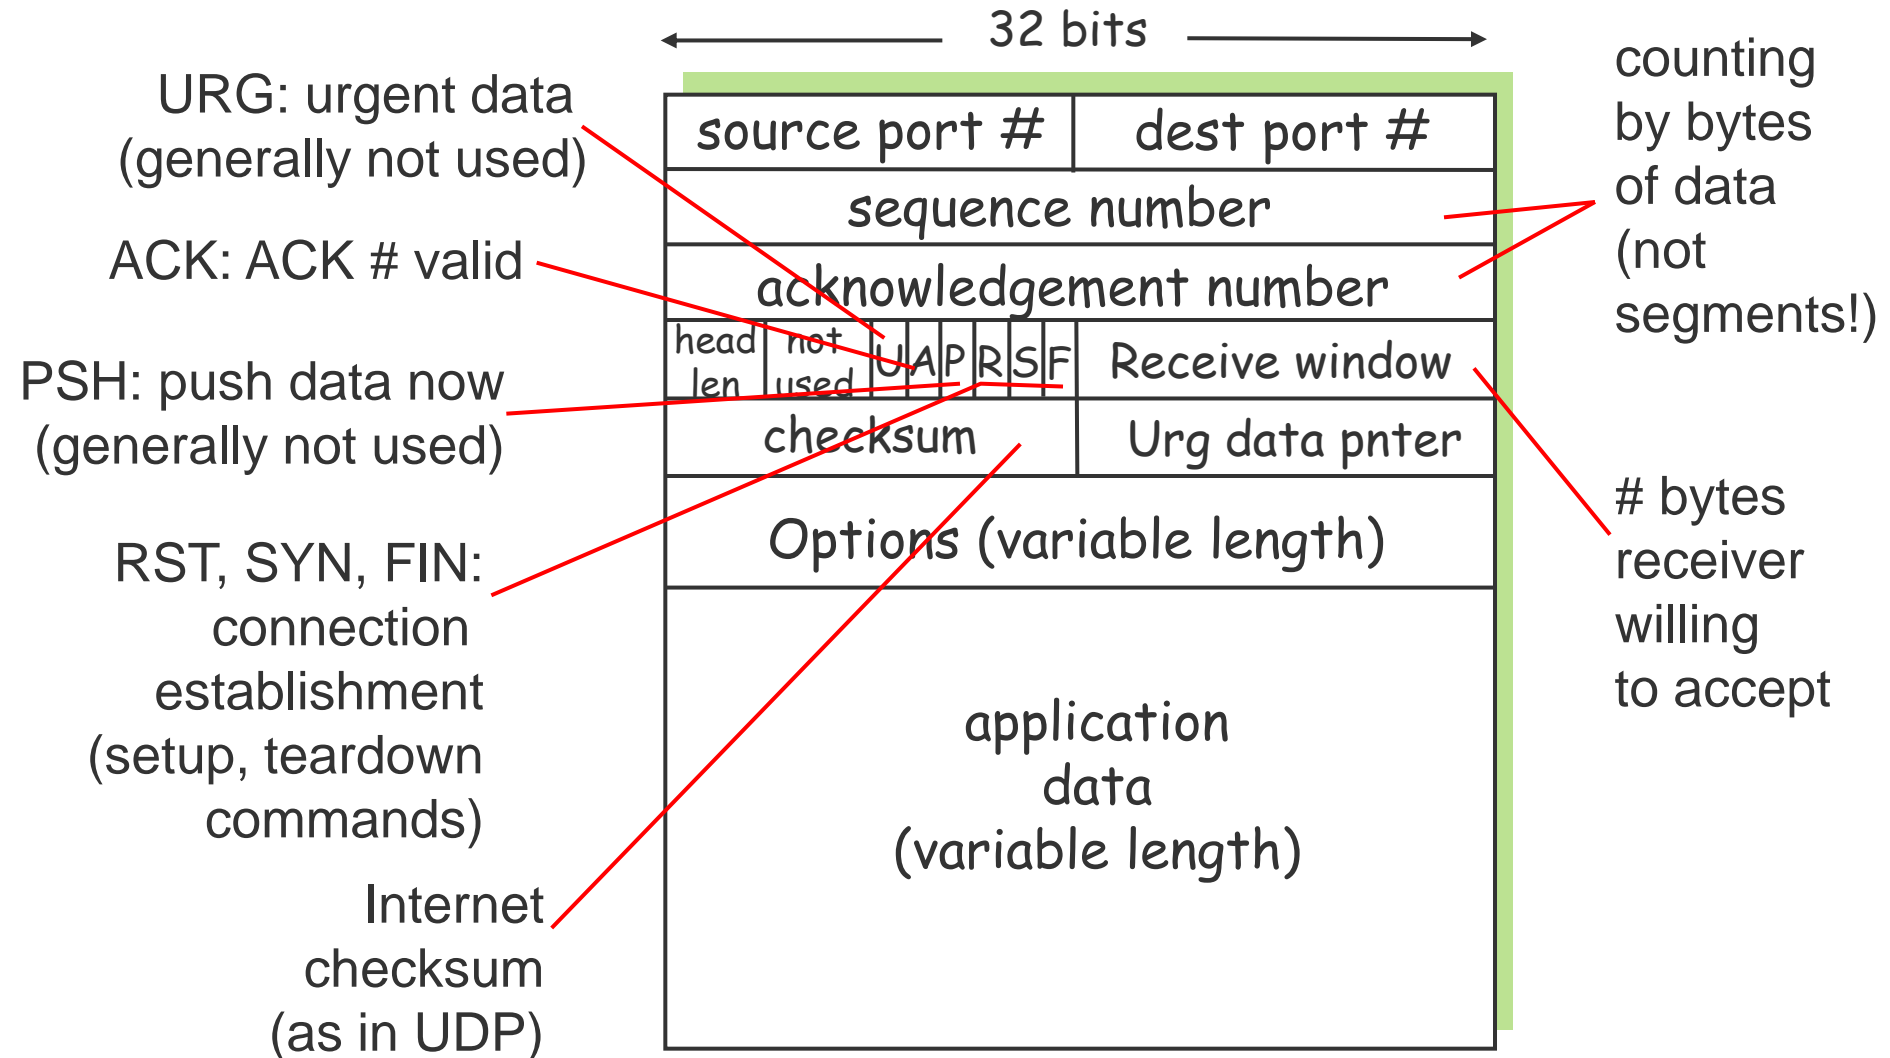
TCP: Connection Flow Control

- Uses variant of sliding windows algorithm
 - window size can be changed

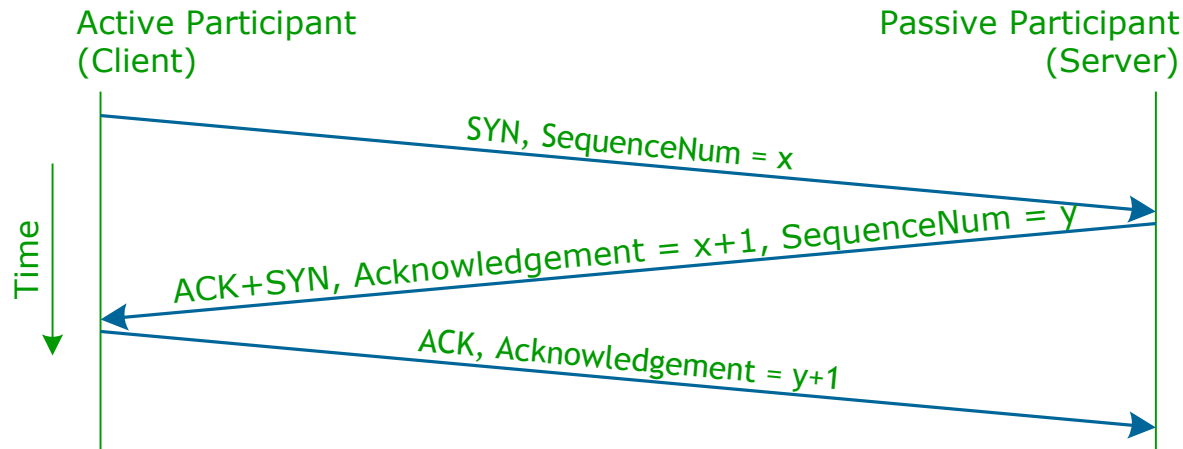


- Can get deadlock

TCP: Segment Format

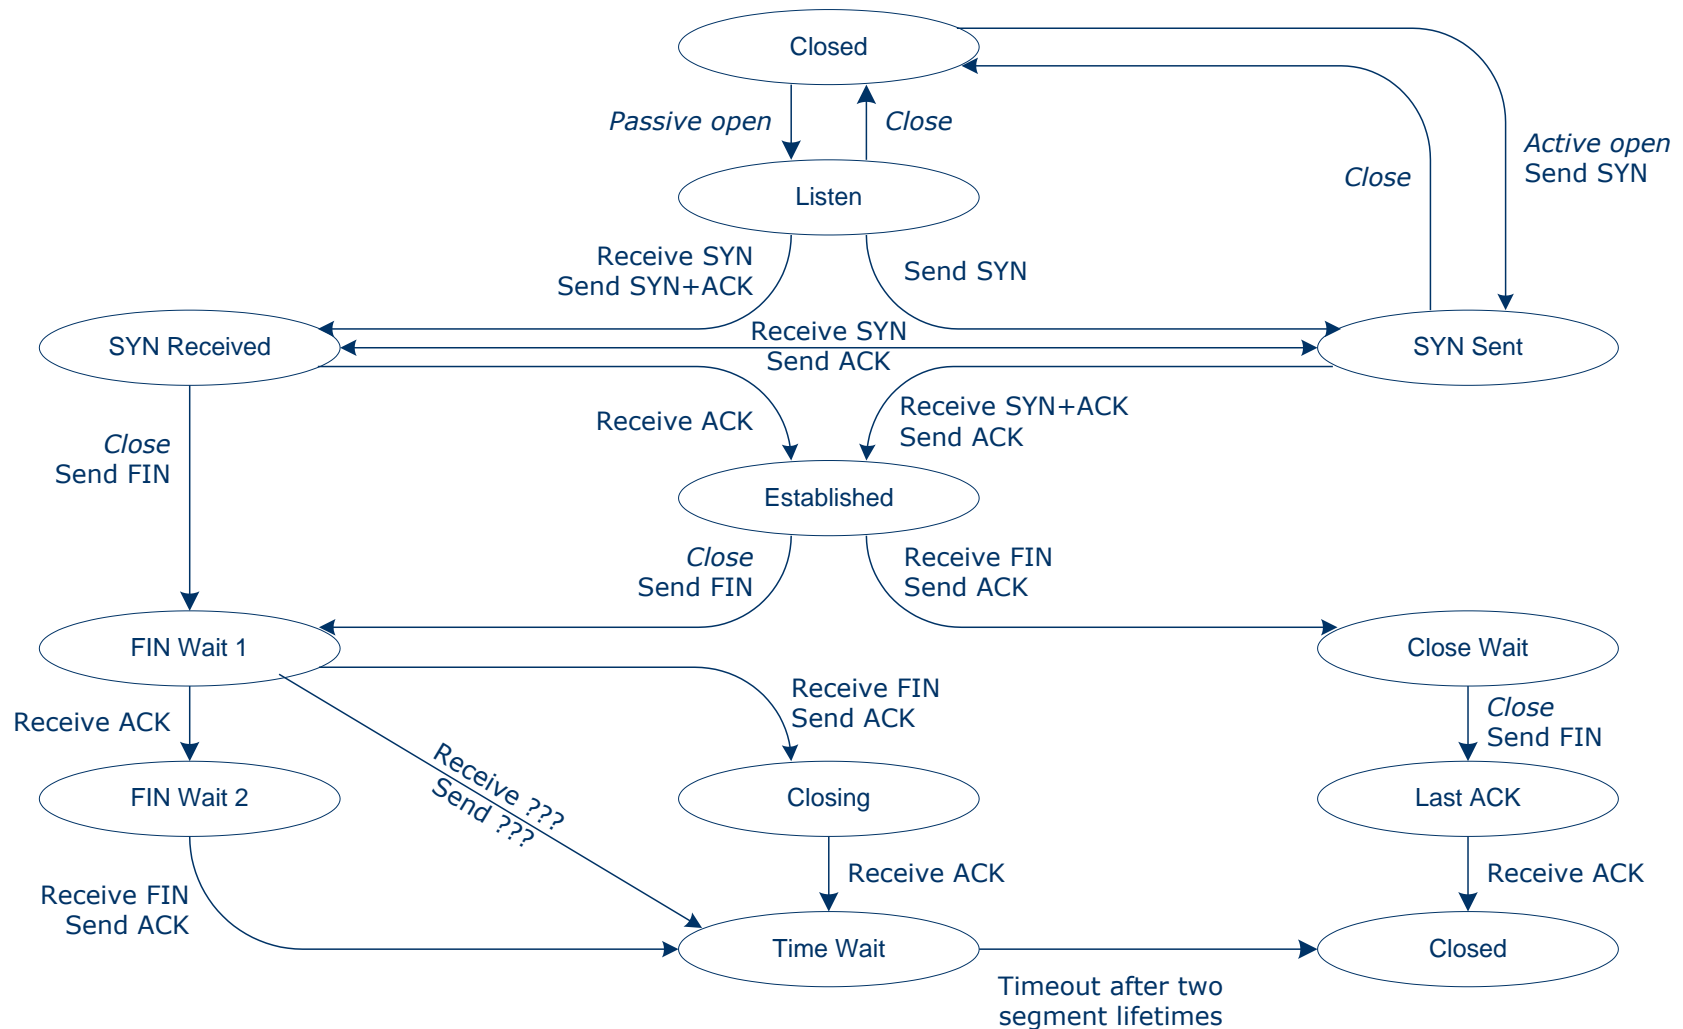


TCP: Establishing Connections



- This exchanges information needed to exchange data
 - initial sequence numbers (ISN) for each end

TCP: State Transition Diagram



TCP: Keeping the Pipe Full

- Capacity of physical connection (100ms RTT)

T1 (1.5Mbps)	18KB	Ethernet (10Mbps)	122KB
T3 (45Mbps)	549KB	FDDI (100Mbps)	1.2MB
STS-3 (155Mbps)	1.8MB	STS-12 (622Mbps)	7.4MB
STS-24 (1.2Gbps)	14.8MB		


- Maximum window of single TCP connection: $2^{16} = 64\text{KB}$
- Maximising use of cables relies on shared use
- Approaching point where:
 - capacity on cable exists but no acknowledgement
- Proposed extension allows:
 - multiplier factor for sequence number and window size

TCP congestion control:

- *goal*: TCP sender should transmit as fast as possible, but without congesting network
 - Q: how to find rate *just* below congestion level
- decentralized: each TCP sender sets its own rate, based on *implicit* feedback:
 - *ACK*: segment received (a good thing!), network not congested, so increase sending rate
 - *lost segment*: assume loss due to congested network, so decrease sending rate

TCP congestion control: bandwidth probing


- “probing for bandwidth”: increase transmission rate on receipt of ACK, until eventually loss occurs, then decrease transmission rate
 - continue to increase on ACK, decrease on loss (since available bandwidth is changing, depending on other connections in network),



- Q: how fast to increase/decrease?
 - details to follow

TCP Fairness


fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K



Why is TCP fair?

Two competing sessions:

- Additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



Transport Summary

- principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- instantiation and implementation in the Internet
 - UDP
 - TCP

Next:

- leaving the network “edge” (application, transport layers)
- into the network “core”