

The University of Manchester

MANCHESTER 1824

Mobile Systems

– real-time computation

COMP28512

Steve Furber & Barry Cheetham

Lecture 8 COMP28512 1

The University of Manchester

MANCHESTER 1824

Real-time computation

- Often computing “as soon as possible” is OK
 - but not for:
 - control, e.g. anti-lock braking
 - streaming media, e.g. mp3 players
 - telephony
 - these all require “real-time” computing
- Hard real-time
 - a missed deadline is a system failure
- Soft real-time
 - a missed deadline is a loss of service quality

Lecture 8 COMP28512 2

The University of Manchester

MANCHESTER 1824

Real-time computation

- Real-time requirements
 - an event should not occur *until* a given time has elapsed
 - an event must take place *before* a given time has elapsed
 - a certain computational activity must take place *within a specified period* of time
 - responding to one event must take *priority* over another event

Lecture 8 COMP28512 3

The University of Manchester

MANCHESTER 1824

Scheduling

- Given
 - the readiness of certain tasks
 - knowledge of how long they take to execute
 - knowledge of their respective deadlines
- Determine which task to execute next

tasks ready now

T1

T2

T3

deadlines

Lecture 8 COMP28512 4

The University of Manchester

MANCHESTER 1824

Responding to external events

- Polling
 - cycle through possible input events repeatedly until an active input is found
 - wastes a lot of power!
- Interrupts
 - an event flag causes the processor to switch attention to handling the event
 - incurs context-switch overheads
- DMA (direct memory access hardware)
 - move the data for processing later

Lecture 8 COMP28512 5

The University of Manchester

MANCHESTER 1824

Buffering

- Handling individual inputs and outputs in real time is near-impossible

input clock

input

process

output clock

output

- what do you do if inputs and outputs must be handled at the same time?

Lecture 8 COMP28512 6

The University of Manchester

Buffering

- Hardware input and output processes can meet timing precisely

- buffers can allow the CPU to process blocks of data
 - more efficient than taking one input at a time

Lecture 8 COMP28512 7

The University of Manchester

Buffering

- In practice we may use DMA to transfer data to buffers in main memory

Lecture 8 COMP28512 8

The University of Manchester

Double-buffers

- Assemble data into blocks
 - work on one block while another is output
 - then switch over to work on the other while the first is output

Lecture 8 COMP28512 9

The University of Manchester

Timing

- Many processor systems incorporate one or more timer/counters
 - run off a system clock
 - can be programmed to interrupt
 - after a certain time
 - at regular intervals
 - can generate output events at precise times
 - provides the processor with a real-time reference

Lecture 8 COMP28512 10

The University of Manchester

Watchdog timers

- What happens if there is a glitch and the system crashes?
- A 'watchdog' timer can be used to interrupt or reset the system
 - normally the software resets the watchdog regularly before it triggers
 - if the software crashes the watchdog will (eventually) trigger
 - functionality is restored, perhaps after a brief loss of real-time performance

Lecture 8 COMP28512 11

The University of Manchester

Real-time communications

- Sending real-time data streams over unreliable, variable delay networks
 - e.g. RTP (Real-time Transport Protocol)
 - what can be done about dropped packets?
 - resend will be too late
 - send packets with
 - sequence number (to see what is missing)
 - timestamp (to get accurate real-time output)
 - no flow control, no error control, no acknowledgement, no retransmission request
 - interpolate to approximate lost data

Lecture 8 COMP28512 12

MANCHESTER 1824

Example: SpiNNaker

Objective:

- to design
 - a chip multiprocessor
 - to model up to a billion spiking neurons in real time
- to investigate
 - biological mechanisms
 - fault-tolerant electronics

Host System Ethernet Link SpiNNaker CMP Asynchronous Interconnect

EPSRC University of Manchester MANCHESTER 1824 ARM Sillistix

Lecture 8 COMP28512 13

MANCHESTER 1824

2D Multi-Chip System

- 2-chip nodes
 - CMP + SDRAM
- Point-to-point links
 - self-timed
 - NoC protocol
- Async packets
 - source key routing
 - key is payload

Lecture 8 COMP28512 14

MANCHESTER 1824

CMP node

Lecture 8 COMP28512 15

MANCHESTER 1824

GALS organization

Lecture 8 COMP28512 16

MANCHESTER 1824

Processor node

Lecture 8 COMP28512 17

MANCHESTER 1824

Event-driven software

Lecture 8 COMP28512 18

The University of Manchester

Real-time kernel

priority level = -1
only one callback
cannot be pre-empted

priority level = 0
can only be pre-empted
by priority -1 callback

priority level > 0
can be pre-empted
by priority ≤ 0 callbacks
scheduled in priority order

Lecture 8 COMP28512 19

The University of Manchester

Power-efficiency

- In a real-time system there are no bonus points to finishing early!
 - when there are no events to process the CPU should go into a low-power mode
 - clocks disabled
 - interrupt wakes CPU up
 - process event
 - go back to sleep
- event-driven model of computation

Lecture 8 COMP28512 20

The University of Manchester

RTOS

- Real-Time Operating Systems
 - class of operating system for real-time systems
 - event-driven priority scheduling
 - switches task only when an event of higher priority needs service
 - task state:
 - running**: currently executing
 - ready**: will run when all higher-priority tasks have completed
 - blocked**: can't run until some enabling event occurs
 - most tasks are blocked most of the time

Lecture 8 COMP28512 21

The University of Manchester

Context switch

- A major issue in real-time computing is *context-switching* overhead
 - saving the state of the old task
 - restoring the state of the new task
 - this can take 10s to 100s of CPU cycles
- Interrupts
 - cause an automatic context switch for interrupt handling
 - often optimised by dedicated hardware resources
 - may cause a pre-emptive task switch

Lecture 8 COMP28512 22

The University of Manchester

Co-routines

- An RTOS may reduce context switch overheads through use of *co-routines*
 - share stack and other process state
 - employ cooperative multitasking
 - instead of pre-emptive multitasking
 - a task hands over control to a higher-priority task at a mutually convenient time
 - must observe various programming restrictions

Lecture 8 COMP28512 23

The University of Manchester

Summary

- Mobile systems often have real-time computation requirements
 - e.g. for streaming media data
- Radio communication is a particular problem
 - because of unavoidable packet loss
- Careful system design minimizes problems
 - buffering, DMA, scheduling, co-routines, ...

Lecture 8 COMP28512 24